

Storing Stuff // TODO

March 21, 2022

Teacher Version

Software Architecture

Evan Hughes & Brae Webb



Figure 1: A map of data storage techniques from Designing Data-Intensive Applications [1].

1 This Week

This week our goal is to:

- explore the various techniques developers use to store data;
- investigate the storage options implementing these techniques on the AWS platform;
- run a small application using docker that requires a database; and
- deploy the application in AWS using Terraform.

2 Databases and Data Models

Unfortunately, to build interesting software we often need to store and use data. The storage of data introduces a number of challenges when designing, creating, and maintaining our software. However, not all data storage techniques are created equal; the choice of data storage model can have a profound impact on our software's complexity and maintainability. In this practical, we want to take a superficial exploration of our island of data storage models. For a more in-depth treatment of data storage models that is outside the scope of this course, see Chapter 2 of the *Designing Data-Intensive Applications* book [1].

For the teacher

Discuss the following different storage technologies and mention some use cases of when you would choose each one. Discuss some popular implementations of each.

Aim for no more than 30 minutes of discussion.

2.1 Relational Storage

Relational databases what have been exposed to the most in your University career — think MySQL, Postgres, Oracle DB, etc. This type of database is good at modelling the real world which is often a highly connected environment.

Some popular offerings are below:

- MySQL/MariaDB [Amazon RDS / Amazon Aurora].
- Postgres [Amazon RDS / Amazon Aurora].

The AWS offerings of these services come in two different types, we have the traditional approach of server capacity (x cores, y ram) and we have a server-less approach. The server-less approach is a more dynamic database that can scale to large amounts of load when needed though at a cost per request.

2.1.1 ORM

Object Relational Mapping (ORM) is a fairly common tool for programmers to use to make developing with databases smoother. One fairly prevalent example of this is SQLAlchemy which is a very widely used database abstraction for python. SQLAlchemy allows us to move to a higher level of abstraction than SQL queries and perform database actions using standard python code.

The benefits of ORMs are the ability to model database objects in our existing programming language instead of having large blocks of SQL text within our source code. The disadvantages come in when we need to do specific SQL work or where the abstractions cost is greater than the benefits.

2.2 Wide-Column Storage

For the teacher

Examples of big apps that depend on this technology is Netflix <https://netflixtechblog.com/netflixs-viewing-data-how-we-know-where-you-are-in-house-of-cards-608dd61077da>.

Wide-Column databases are a form of NoSQL or non-relational data stores. In these data stores the data model design is focused more on having efficient queries at the cost of data duplication. A warning to the reader that these models are not flexible after creation, it is much easier to answer a new use case in a relational model.

- Apache Cassandra [Amazon Keyspaces for Cassandra].
- Apache HBase.

2.3 Key-Value Storage

Key-Value stores are very popular for cache or remote config use cases, some of the most notable are Redis and Memcached. These stores allow efficient lookup of values via keys and are usually stored in-memory.

- Redis [Amazon ElastiCache for Redis].
- Memcached [Amazon ElastiCache for Memcached].
- Amazon DynamoDB.
- Amazon MemoryDB for Redis.

2.4 Time Series Storage

For the teacher

Something to mention here is that relations are usually not utilised between tables in time series databases.

Time series databases are highly focused storage which is tailored to retrieving results by timestamp ranges. Many implementations also take advantage of the data model to allow efficient rollover of data and partitioning. One of the most popular time series databases is Prometheus which is used to store monitoring metrics.

- Amazon Timestream.
- TimescaleDB (Postgres + Addon).
- Prometheus.
- InfluxDB.

2.5 Document Storage

Document databases are a subset of NoSQL databases with a focus on a flexible data model. MongoDB for instance allows the user to store JSON documents and perform queries on those documents. One advantage of document databases is that they match a programmers existing mental model of storing data in formats such as JSON.

- MongoDB.
- Apache CouchDB.
- Amazon DocumentDB.
- Amazon DynamoDB.

2.6 Graph Storage

For the teacher

If you havnt experienced graph databases, a good usecase is “recommendation systems”, which use the connected nature of items to figure out what to suggest to a person. Another example is the <https://neo4j.com/blog/analyzing-panama-papers-neo4j/> Panama Papers.

Graph Databases are relational storage with a few enhancements to allow fast neighbour look-ups. These databases also allow the implementation of graph algorithms to query data.

- Amazon Neptune.
- Neo4J.
- Janus Graph.

3 Deploying an Application with Storage

Thus far in the course we have introduced docker in small packages which do not communicate (much) with the outside world. Today we will be using it to run an application which needs to talk to a database.

1. First we will deploy the application and database locally.
2. Then we will deploy the database on AWS infrastructure and configure our application to talk to the remote database.
3. Finally, we will deploy the application and database on AWS infrastructure.

My Todo List

Complete CSSE6400 Prac 1	+
Complete CSSE6400 Prac 2	
Complete CSSE6400 Prac 3	
Complete CSSE6400 Prac 4	
Joined the CSSE6400 Slack	
Attended Lecture 1 of CSSE6400	
Attended Lecture 2 of CSSE6400	
Attended Lecture 3 of CSSE6400	
Attended Lecture 4 of CSSE6400	
Attended Braes tutorial	

Previous

1

2

3

Next

Figure 2: Sample Todo App made for the course

Info

You will need to have docker and docker-compose installed for this practical. Installation will depend on your operating system.

- Docker compose: <https://docs.docker.com/compose/install/>
- Docker engine: <https://docs.docker.com/get-docker/>

We also recommend installing the vscode docker plugin or the equivalent tools in IntelliJ IDEs.

Aside

We have only seen individual Docker containers running thus far. As we start to build distributed systems, we will want to use multiple Docker containers which talk to each other. Docker compose is a container orchestration tool that allows us to configure multiple containers to run and setup paths for them to talk to each other.

For the teacher

Wait for students to get docker-compose installed, they should have docker from their tutorials but some may be missing it.

Warning

For terminal examples in this section, lines that begin with a \$ indicate a line which you should type while the other lines are example output that you should expect. Not all of the output is captured in the examples to save on space.

4 Deploying Locally

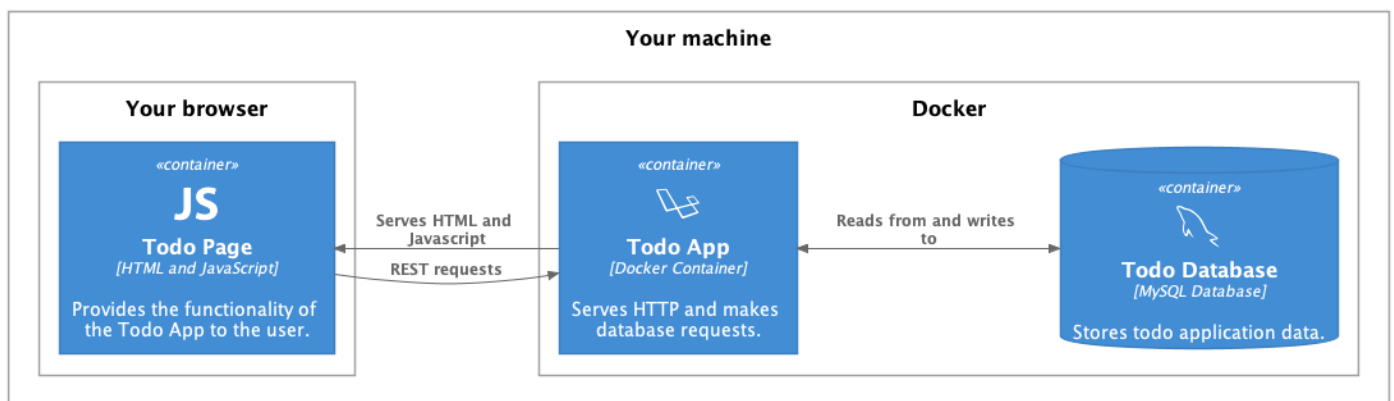


Figure 3: Local deployment diagram

For the teacher

Mention that the dockerfile exists but no need to get the repo. We will not be building the container ourselves. Instead use one that is published on the github, shown further down in the docker-compose.

We will be using a container that is pre-built from the Dockerfile below. You should browse the contents of this container to gain more exposure to Dockerfile although you aren't required to use the file for this task — we will use an image already built from this Dockerfile.

```
» cat Dockerfile
```

```
1 FROM ubuntu:21.10
2 RUN apt-get update \
3     && DEBIAN_FRONTEND=noninteractive apt install -y \
4         php \
5         php-mysql \
6         php-xml \
7         php-curl \
8         curl \
```

```

9         git \
10        unzip
11 RUN curl -sS https://getcomposer.org/installer | php -- --install-dir=/usr/local/bin
    --filename=composer
12 COPY . /app
13 WORKDIR /app
14 RUN composer install
15 CMD ["php", "artisan", "serve", "--host=0.0.0.0"]

```

Our first task is to have a local instance of the Todo App including the database. To get started we need to make a new directory for our work and create a Docker compose file.

```

$ mkdir prac4 && cd prac4
$ touch docker-compose.yml

```

Docker Compose is a small helper utility that allows us to more easily run docker applications without needing to remember a lot of command line parameters. Instead we define how we want our docker container to run through a YAML config file. Insert the following into your docker-compose.yml file.

```

» cat docker-compose.yml
1 version: '3.3'
2 services:
3   backend:
4     image: ghcr.io/csse6400/todo-app:combined-latest
5     ports:
6       - '8000:8000'
7     environment:
8       APP_ENV: 'local'
9       APP_KEY: 'base64:8PQEPYGlTm1t3aqWmlAw/ZPwCiIFvdXDBjk3mhsom/A='
10      APP_DEBUG: 'true'
11      LOG_LEVEL: 'debug'

```

For the teacher

Feel free to show students dockerhub and where on github this container is stored. URL for this container is here <https://github.com/CSSE6400/todo-app/pkgs/container/todo-app>

A few things to point out in the file. We have defined a single service (container) called backend which uses the pre-built docker image from ghcr.io/csse6400/todo-app with the tag of combined-latest. We then have exposed this onto our machine on port 8000 and have passed a few environment variables required for the application.

`docker-compose up` will use our docker compose configuration to run any containers we have specified.

```
$ docker-compose up
Creating network "p1_default" with the default driver
Creating p1_backend_1 ... done
Attaching to p1_backend_1
backend_1 | Starting Laravel development server: http://0.0.0.0:8000
backend_1 | [Sun Mar 20 07:56:23 2022] PHP 8.0.8 Development Server (http
: //0.0.0.0:8000) started
```

Now we head to our browser and go to <http://127.0.0.1:8000>, you should be presented with the following screen.

My Todo List

Previous

1

Next



The server had a problem

When we see unhelpful error messages such as this, we should investigate the web console in our browser. The web console will tell us which network calls are being made and which calls are causing the error message. Follow along with your tutor to find that the offending network call is to <http://127.0.0.1:8000/api/v1/todo>.

Once we open that address we have a clearer idea of what has gone wrong. The page you should see is shown in Figure 4 and the reason is that our backend cannot connect to a database (because we have yet to make one).




Figure 4: The expected error page when accessing <http://127.0.0.1:8000/api/v1/todo>.

To fix this lets add a popular relation database, MySQL, to our docker-compose file. Edit your docker compose file to match as shown below.

```
» cat main.tf
1 version: '3.3'
2 services:
3   db:
4     image: mysql:8-debian
5     environment:
6       MYSQL_DATABASE: 'todoapp'
7       MYSQL_USER: 'todoapp'
8       MYSQL_PASSWORD: 'password'
9       MYSQL_ROOT_PASSWORD: 'password'
10    ports:
11      - '3306:3306'
13 backend:
```

```

14 image: ghcr.io/csse6400/todo-app:combined-latest
15 depends_on:
16   - db
17 ports:
18   - '8000:8000'
19 environment:
20   APP_ENV: 'local'
21   APP_KEY: 'base64:8PQEPYGlTm1t3aqWmlAw/ZPwCiIFvdXDBjk3mhsom/A='
22   APP_DEBUG: 'true'
23   LOG_LEVEL: 'debug'
24   DB_CONNECTION: 'mysql'
25   DB_HOST: 'db'
26   DB_PORT: '3306'
27   DB_DATABASE: 'todoapp'
28   DB_USERNAME: 'todoapp'
29   DB_PASSWORD: 'password'

```

Now we have two services (containers) for our app and we have added a few more environment variables for our backend to know how to connect to the database. The `DB_HOST` variable uses a feature of docker compose where you can refer to other services by their name rather than a traditional IP address. This makes it easy for us to setup communication between these two services.

From the same shell let's re-run our containers, you may need to CTRL+C to stop the current running containers. Once they have shutdown, run the up command again.

```

$ docker-compose up
Starting p2_db_1 ... done
Starting p2_backend_1 ... done
Attaching to p2_db_1, p2_backend_1
db_1 | 2022-03-20 08:11:55+00:00 [Note] [Entrypoint]: Entrypoint ....
db_1 | 2022-03-20 08:11:55+00:00 [Note] [Entrypoint]: Switching t....
db_1 | 2022-03-20 08:11:55+00:00 [Note] [Entrypoint]: Entrypoint ....
db_1 | 2022-03-20T08:11:55.438996Z 0 [System] [MY-010116] [Server....
db_1 | 2022-03-20T08:11:55.445261Z 1 [System] [MY-013576] [InnoDB....
backend_1 | Starting Laravel development server: http://0.0.0.0:8000
db_1 | 2022-03-20T08:11:55.535803Z 1 [System] [MY-013577] [InnoDB....
db_1 | 2022-03-20T08:11:55.673757Z 0 [Warning] [MY-010068] [Serve....
db_1 | 2022-03-20T08:11:55.673784Z 0 [System] [MY-013602] [Server....
db_1 | 2022-03-20T08:11:55.674810Z 0 [Warning] [MY-011810] [Serve....
db_1 | 2022-03-20T08:11:55.684729Z 0 [System] [MY-010931] [Server....
db_1 | 2022-03-20T08:11:55.684756Z 0 [System] [MY-011323] [Server....
backend_1 | [Sun Mar 20 08:11:55 2022] PHP 8.0.8 Development Serv....

```

Now when we go to <http://127.0.0.1:8000/api/v1/todo> we see a different error message, as shown in Figure 5. This error is complaining that we have a database that we can connect to but the todos table doesn't exist.

To populate the database our application comes with database migration files. One way would be to click the "RUN MIGRATIONS" button shown on the error page but we also want to pre-populate our database with some dummy data as well.




Figure 5: The expected error page when accessing <http://127.0.0.1:8000/api/v1/todo> after creating a database.

To do this we are going to jump into the running container and execute the migrations ourselves, using `docker-compose exec`. Start by opening a new terminal so that we can leave the docker containers running. In this new terminal navigate to **the directory where `docker-compose.yml` lives** and run the following command:

```
$ docker-compose exec backend php artisan migrate:fresh --seed
Dropped all tables successfully.
Migration table created successfully.
Migrating: 2022_03_19_041557_create_todos_table
Migrated: 2022_03_19_041557_create_todos_table (7.55ms)
Seeding: Database\Seeders\TodoSeeder
Seeded: Database\Seeders\TodoSeeder (6.56ms)
Database seeding completed successfully.
```

Now with this run we can check back at our web app (at <http://127.0.0.1:8000>) and you should see a fully functional todo app.

My Todo List

Complete CSSE6400 Prac 1	+
Complete CSSE6400 Prac 2	
Complete CSSE6400 Prac 3	
Complete CSSE6400 Prac 4	
Joined the CSSE6400 Slack	
Attended Lecture 1 of CSSE6400	
Attended Lecture 2 of CSSE6400	
Attended Lecture 3 of CSSE6400	
Attended Lecture 4 of CSSE6400	
Attended Braes tutorial	

[Previous](#) [1](#) [2](#) [3](#) [Next](#)

Info

The database migrations are performed by Laravel which is a popular PHP web framework. In the database it has created a table to keep track of which migrations have already been run so that it will skip them latter. To enable this functionality you will have to edit `migrate:fresh` to just `migrate`.

For a production instance you would also typically remove the `--seed` parameter as you do not want to insert dummy data into your production database.

4.1 Exercise: Migrations at startup

So far when we run this application we have to perform the database migrations manually. To help us get up and running we are going to make a small modification to pre-run the migrations when are web app starts. First we need to have a look at how the container is set to launch by default. In the Dockerfile attached at the start of the practical we see that we have defined the command to run on the last line with the `CMD` directive.

```
» cat Dockerfile
```

```
1 FROM ubuntu:21.10
2 ...
3 ...
4 ...
5 CMD ["php", "artisan", "serve", "--host=0.0.0.0"]
```

Info

When working with docker it can get confusing around the networking aspects. In this application I have specified that the server must listen on all network interfaces (0.0.0.0). Without this flag the default is 127.0.0.1 which even though its the localhost the forwarded traffic through the docker container would never reach it.

This command launches the laravel development server and listens on all interfaces on the host. We are going to override this in our docker-compose file so that we run the migrations then start the server. Add the following line to the docker-compose.yml that you have been developing during the practical.

```
1 command: sh -c "sleep 30 && php artisan migrate:refresh --seed && php artisan serve
    --host=0.0.0.0"
```

This new command does the following:

- Waits for the database to be ready in a simple way.
- Runs the migrations and seeds the database, as we have seen earlier.
- Starts the development server as the container originally did.

Example: condensed version of the goal docker-compose.yml attached below.

```
» cat docker-compose.yml
1 version: '3.3'
2 services:
3   db:
4     ...
5
6   backend:
7     ...
8     environment:
9       ...
10    command: sh -c "sleep 10 && php artisan migrate:refresh --seed && php artisan
    serve --host=0.0.0.0"
```

Now when we launch the docker-compose we can see that our migrations were run in the output.

```
$ docker-compose up
...
backend_1 | Rolling back: 2022_03_19_041557_create_todos_table
backend_1 | Rolled back: 2022_03_19_041557_create_todos_table (8.28ms)
backend_1 | Migrating: 2022_03_19_041557_create_todos_table
backend_1 | Migrated: 2022_03_19_041557_create_todos_table (11.55ms)
backend_1 | Seeding: Database\Seeders\TodoSeeder
backend_1 | Seeded: Database\Seeders\TodoSeeder (44.77ms)
backend_1 | Database seeding completed successfully.
backend_1 | Starting Laravel development server: http://0.0.0.0:8000
backend_1 | [Sun Mar 20 12:08:41 2022] PHP 8.0.8 Development Server (http
://0.0.0.0:8000) started
```

We can also bake this into the container by extending the original, it is fairly common to see projects in the wild that run a init script when the container launches. An exercise left for the reader is to build upon the provided docker container by including an init script.

5 Deploying a Database in AWS

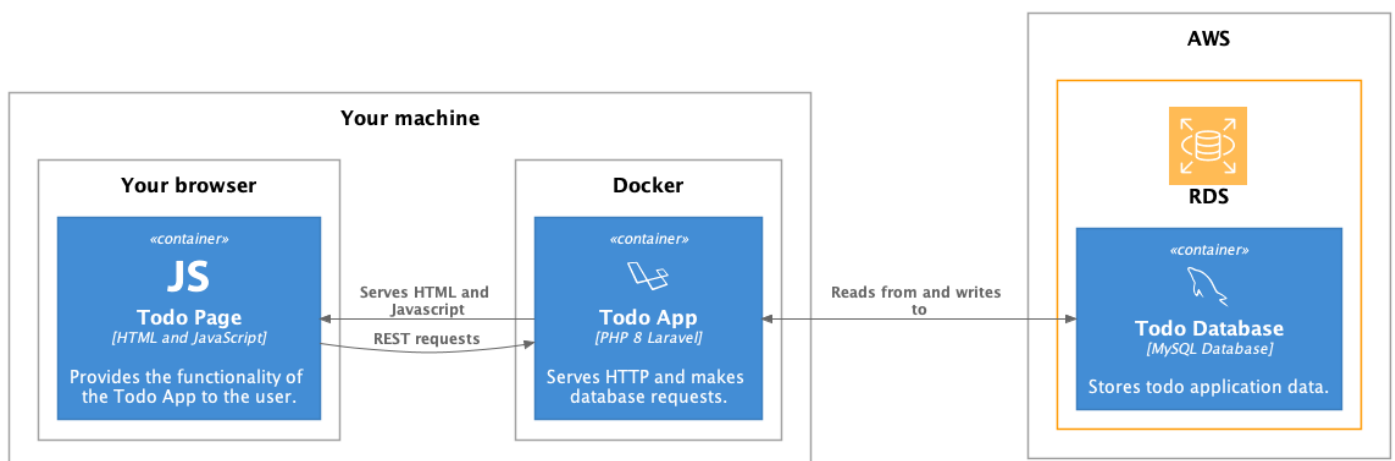



Figure 6: Remote database deployment diagram

Now we have a locally running Todo App let's move to AWS, start up your Learner Lab environment now. This is the last time we will heavily use the AWS user interface in the practicals. If you already feel confident in the AWS environment skip to Section 6 for the terraform setup.

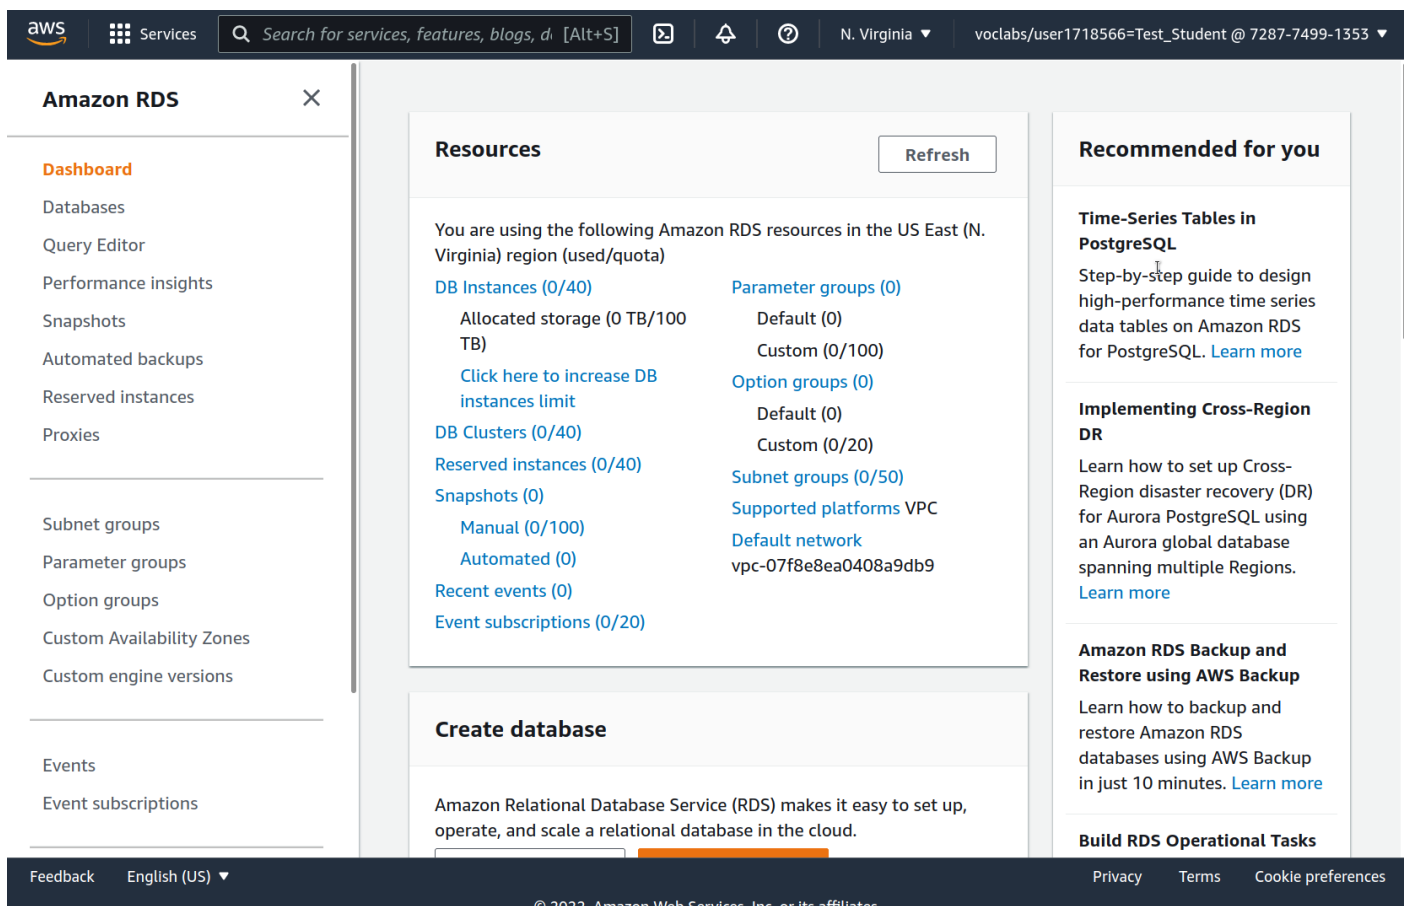
For the teacher

Give students time to start the labs, could take up to 10 minutes.

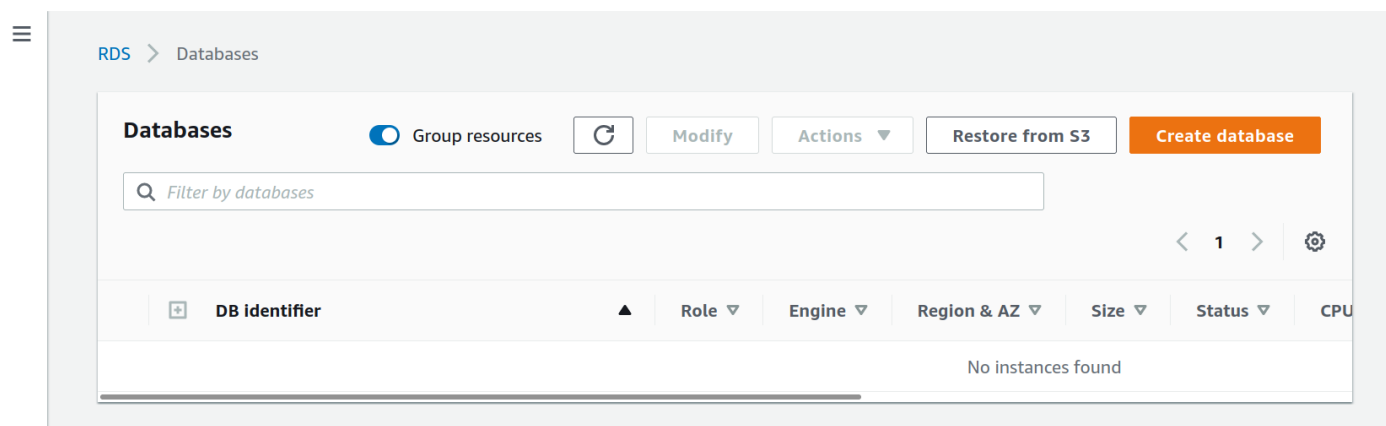
To get started let's jump into the lab environment and have a look at AWS RDS which is an AWS managed database service. To get to the RDS service either search it or browse Services -> Database -> RDS as shown below.



Now we are in the management page for all our database instances, for today we just want to get a small instance running to explore the service. Head to “DB Instances (0/40)”.



This page should appear familiar as it's very similar to the AWS EC2 instance page. Let us create a new database by hitting the "Create Database" button.



Warning

In the next section we cannot use the Easy Create method as it tries to create a IAM account which is not allowed in the labs. Going forward we would typically do this using Terraform so we can easily avoid these restrictions.

For the teacher

Feel free to talk about the other offerings here, but make sure to flame Oracle and Microsoft SQL Server. A good thing to point out is the Amazon Aurora which is the serverless version of RDS.

We will be creating a standard database so select standard and MySQL. We will use version 8 to match the local version.

Choose a database creation method [Info](#)

☒ Standard create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

☐ Easy create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

☐ Amazon Aurora



☒ MySQL



☐ MariaDB



☐ PostgreSQL



☐ Oracle



☐ Microsoft SQL Server



Edition

☒ MySQL Community



Known issues/limitations

Review the [Known issues/limitations](#) [link](#) to learn about potential compatibility issues with specific database versions.

Version

MySQL 8.0.27



For today we are going to use “Free Tier” but in the future, you may wish to explore the different deployment options. Please peruse the available different options.

For the teacher

Walk through what Multi-AZ means aka Multiple Availability Zones.

Templates

Choose a sample template to meet your use case.



Production

Use defaults for high availability and fast, consistent performance.



Dev/Test

This instance is intended for development use outside of a production environment.



Free tier

Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS.

[Info](#)

Availability and durability

Deployment options [Info](#)

The deployment options below are limited to those supported by the engine you selected above.



Single DB instance (not supported for Multi-AZ DB cluster snapshot)

Creates a single DB instance with no standby DB instances.



Multi-AZ DB instance (not supported for Multi-AZ DB cluster snapshot)

Creates a primary DB instance and a standby DB instance in a different AZ. Provides high availability and data redundancy, but the standby DB instance doesn't support connections for read workloads.



Multi-AZ DB Cluster - new

Creates a DB cluster with a primary DB instance and two readable standby DB instances, with each DB instance in a different Availability Zone (AZ). Provides high availability, data redundancy and increases capacity to serve read workloads.

Now we need to name our database and create credentials to connect via. Please enter a reasonable password and keep this aside for later. We will need it for our local docker-compose file.

Settings

DB instance identifier [Info](#)

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in the current AWS Region.

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1 to 60 alphanumeric characters or hyphens. First character must be a letter. Can't contain two consecutive hyphens. Can't end with a hyphen.

▼ Credentials Settings

Master username [Info](#)

Type a login ID for the master user of your DB instance.

1 to 16 alphanumeric characters. First character must be a letter.

☐ Auto generate a password

Amazon RDS can generate a password for you, or you can specify your own password.

Master password [Info](#)

Constraints: At least 8 printable ASCII characters. Can't contain any of the following: / (slash), ' (single quote), " (double quote) and @ (at sign).

Confirm password [Info](#)

We will use the default class type, t2.micro, which should be sufficient for this practical.

For the teacher

May want to mention that burstable is not recommended for consistently used databases. Usually DBs are memory focused and thus the standard or memory optimised are used.

DB instance class

DB instance class [Info](#)

☐ Standard classes (includes m classes)

☐ Memory optimized classes (includes r and x classes)

☒ Burstable classes (includes t classes)

1 vCPUs 1 GiB RAM Not EBS Optimized

☐ Include previous generation classes

For storage we will leave all the default options.

Storage

Storage type [Info](#)

General Purpose SSD (gp2)
Baseline performance determined by volume size

Allocated storage

20

GiB

(Minimum: 20 GiB. Maximum: 16,384 GiB) Higher allocated storage **may improve** IOPS performance.

You might see better baseline performance with your selected volume size by specifying General Purpose SSD storage. [Learn more about using Provisioned IOPS storage for consistent performance.](#)

Storage autoscaling [Info](#)

Provides dynamic scaling support for your database's storage based on your application's needs.

☒

Enable storage autoscaling

Enabling this feature will allow the storage to increase once the specified threshold is exceeded.

Maximum storage threshold [Info](#)

Charges will apply when your database autoscales to the specified threshold

1000

GiB

Minimum: 21 GiB. Maximum: 16,384 GiB

In connectivity we need to make sure our instance is publicly available. Usually you don't want to expose your databases publicly and, would instead, have a web server sitting in-front. But for today we will be running that web server locally so for convenience we need public access.

When selecting public access as yes we have to create a new Security Group, give this Security Group a sensible name.

20

Connectivity



Virtual private cloud (VPC) [Info](#)

VPC that defines the virtual networking environment for this DB instance.

Default VPC (vpc-07f8e8ea0408a9db9) ▼

Only VPCs with a corresponding DB subnet group are listed.

After a database is created, you can't change its VPC.

Subnet group [Info](#)

DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

default-vpc-07f8e8ea0408a9db9 ▼

Public access [Info](#)

☒ Yes

Amazon EC2 instances and devices outside the VPC can connect to your database. Choose one or more VPC security groups that specify which EC2 instances and devices inside the VPC can connect to the database.

☐ No

RDS will not assign a public IP address to the database. Only Amazon EC2 instances and devices inside the VPC can connect to your database.

VPC security group

Choose a VPC security group to allow access to your database. Ensure that the security group rules allow the appropriate incoming traffic.



Choose existing

Choose existing VPC security groups



Create new

Create new VPC security group

New VPC security group name

todoapp-manual

Availability Zone [Info](#)

No preference ▼

▼ Additional configuration

Database port [Info](#)

TCP/IP port that the database will use for application connections.

3306



We will leave the authentication as password based but we need to expand the “Additional configuration”. Fill in the “Initial Database Name” section to be “todoapp”, this is similar to what we had in the Docker Compose.

For the teacher

The other options here are to do with the parameters used to start the database, it is uncommon to have to change these but this is where any settings you would pass in via cli to the db would be set.

Database authentication

Database authentication options [Info](#)

- ☒ **Password authentication**
Authenticates using database passwords.
- ☐ **Password and IAM database authentication**
Authenticates using the database password and user credentials through AWS IAM users and roles.
- ☐ **Password and Kerberos authentication**
Choose a directory in which you want to allow authorized users to authenticate with this DB instance using Kerberos Authentication.

▼ Additional configuration

Database options, backup enabled, backtrack disabled, Enhanced Monitoring disabled, maintenance, CloudWatch Logs, delete protection disabled.

Database options

Initial database name [Info](#)

If you do not specify a database name, Amazon RDS does not create a database.

DB parameter group [Info](#)

Option group [Info](#)

Now we can click create which will take some time.


Estimated monthly costs

The Amazon RDS Free Tier is available to you for 12 months. Each calendar month, the free tier will allow you to use the Amazon RDS resources listed below for free:

- 750 hrs of Amazon RDS in a Single-AZ db.t2.micro Instance.
- 20 GB of General Purpose Storage (SSD).
- 20 GB for automated backup storage and any user-initiated DB Snapshots.

[Learn more about AWS Free Tier.](#)

When your free usage expires or if your application use exceeds the free usage tiers, you simply pay standard, pay-as-you-go service rates as described in the [Amazon RDS Pricing page.](#)






 You are responsible for ensuring that you have all of the necessary rights for any third-party products or services that you use with AWS services.


Cancel

Create database

Depending on your database it may take 10 to 30 minutes to create, the larger and more complicated the setup the longer it usually takes. The database will also do a initial backup when its created.

RDS > Databases

Databases ☒ Group resources     

 DB identifier

Role



Engine

Region & AZ

Size

Status

CPU

 todoapp-manual	Instance	MySQL Community	us-east-1f	db.t2.micro	 Backing-up	<div><div></div><div>100%</div></div>
--	----------	-----------------	------------	-------------	--	---------------------------------------

When the database has finished being created you can select it to view the configuration and details. In this menu we also see the endpoint address which we will need to copy into our docker compose file.

Successfully created database [todoapp-manual](#)

View connection details

✕

[RDS](#) > [Databases](#) > todoapp-manual

todoapp-manual

Modify

Actions ▼

Summary

DB identifier todoapp-manual	CPU <div><div></div></div> 10.83%	Status Available	Class db.t2.micro
Role Instance	Current activity <div><div></div></div> 0 Connections	Engine MySQL Community	Region & AZ us-east-1f

Connectivity & security

Monitoring

Logs & events

Configuration

Maintenance & backups

Tags

Connectivity & security

<div>Endpoint & port</div> <div>Endpoint todoapp-manual.cwvf1cdgoxax.us-east-1.rds.amazonaws.com</div> <div>Port 3306</div>	<div>Networking</div> <div>Availability Zone us-east-1f</div> <div>VPC vpc-07f8e8ea0408a9db9</div> <div>Subnet group default-vpc-07f8e8ea0408a9db9</div> <div>Subnets subnet-0556db549e22800f1 subnet-0da793726639bd6d8 subnet-091ee1d302ae831a9 subnet-0b932c4d6a4154b2a subnet-0416084227ea643b4 subnet-05d92ccc16a62294f </div>	<div>Security</div> <div>VPC security groups todoapp-manual2 (sg-0cc1a6ba52b85e23c) Active </div> <div>Publicly accessible Yes</div> <div>Certificate authority rds-ca-2019</div> <div>Certificate authority date August 23, 2024, 03:08 (UTC±3:08)</div>
---	--	---

So with the database all ready to go, let's get started on running our app. First thing to do is to stop any running instances you have and then delete the database service from the docker compose. Now we want to edit the DB_HOST and DB_PASSWORD to match what we set in AWS. Below is an example of my configuration.

```

1  version: '3.3'
2  services:
3    backend:
4      image: ghcr.io/csse6400/todo-app:combined-latest
5      ports:
6        - '8000:8000'
7      environment:
8        APP_ENV: 'local'
9        APP_KEY: 'base64:8PQEPYGI7m1t3aqWmlAw/ZPwCiIFvdXDBjk3mhsom/A='
10       APP_DEBUG: 'true'
11       LOG_LEVEL: 'debug'
12       DB_CONNECTION: 'mysql'
13       DB_HOST: 'todoapp-manual.cwf1cdgoxzax.us-east-1.rds.amazonaws.com'
14       DB_PORT: '3306'

```



```

15 DB_DATABASE: 'todoapp'
16 DB_USERNAME: 'todoapp'
17 DB_PASSWORD: 'MyVerySecurePassword'
18 command: sh -c "sleep 10 && php artisan migrate:refresh --seed && php artisan
    serve --host=0.0.0.0"

```

Now we can run `$ docker-compose up`. This will take a lot longer than before, the reason behind this is that by default our lab is running in the US and it takes a lot longer to communicate to a DB half way across the world.

```

$ docker-compose up
Creating network "remote_default" with the default driver
Creating remote_backend_1 ... done
Attaching to remote_backend_1
backend_1 | Migration table not found.
backend_1 | Migration table created successfully.
backend_1 | Migrating: 2022_03_19_041557_create_todos_table
backend_1 | Migrated: 2022_03_19_041557_create_todos_table (612.48ms)
backend_1 | Seeding: Database\Seeders\TodoSeeder
backend_1 | Seeded: Database\Seeders\TodoSeeder (13,930.36ms)
backend_1 | Database seeding completed successfully.
backend_1 | Starting Laravel development server: http://0.0.0.0:8000
backend_1 | [Mon Mar 21 04:45:34 2022] PHP 8.0.8 Development Server (http
    ://0.0.0.0:8000) started

```

We can now also browse the app as before but with just a lot more lag.

6 Deploying in AWS




Figure 7: Remote deployment diagram

In this section we will deploy the application and database in AWS. Rather than manually configuring a database as we did in the previous section, we will deploy everything with Terraform.

6.1 Authentication

As we did last week, specify the AWS provider we will use and download the Learner Lab credentials into a credentials file. Once you have setup your `main.tf` and `credentials` files, run `terraform init`.

```
» cat main.tf
1 terraform {
2   required_providers {
3     aws = {
4       source = "hashicorp/aws"
5       version = "~> 3.0"
6     }
7   }
8 }
10 provider "aws" {
11   region = "us-east-1"
12   shared_credentials_file = "./credentials"
13 }
```

6.2 RDS Database

Now would be a good time to browse the documentation for the RDS database in Terraform: https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/db_instance. Using our manual configuration, we can come up with a resource with the appropriate parameters as below:

```
» cat main.tf
1 locals {
2   password = "foobarbaz" # this is bad
3 }
5 resource "aws_db_instance" "todoapp-database" {
6   allocated_storage = 20
7   max_allocated_storage = 1000
8   engine = "mysql"
9   engine_version = "8.0.27"
10  instance_class = "db.t2.micro"
11  name = "todoapp"
12  username = "todoapp"
13  password = local.password
14  parameter_group_name = "default.mysql8.0"
15  skip_final_snapshot = true
16  vpc_security_group_ids = [aws_security_group.todoapp-database.id]
17  publicly_accessible = true
19  tags = {
20    Name = "todoapp-database"
  }
```

```
21 }
22 }
```

Remember to create an appropriate security group as we did through the user interface.

```
» cat main.tf

1 resource "aws_security_group" "todoapp-database" {
2   name = "todoapp-database"
3   description = "Allow inbound MySQL traffic"
4
5   ingress {
6     from_port = 3306
7     to_port = 3306
8     protocol = "tcp"
9     cidr_blocks = ["0.0.0.0/0"]
10  }
11
12  egress {
13    from_port = 0
14    to_port = 0
15    protocol = "-1"
16    cidr_blocks = ["0.0.0.0/0"]
17    ipv6_cidr_blocks = ["::/0"]
18  }
19
20  tags = {
21    Name = "todoapp-database"
22  }
23 }
```

6.3 Container on AWS

As we mentioned in the Infrastructure as Code notes [2], in this course we will use Docker to configure machines and Terraform to configure infrastructure. AWS has the ability to deploy Docker containers using a service known as Elastic Container Service (ECS). Unfortunately, the AWS Learner Labs provided by AWS do not support ECS.

To resolve this issue, we have created a Terraform module which allows us to deploy Docker images on EC2 instances and abstract over the underlying implementation. The documentation and source for this Terraform module is available on Github: <https://github.com/CSSE6400/terraform/tree/main/container>.

Using the documentation of the module, combined with the environment variables we know our backend requires based on the `docker-compose.yml` file, we can develop a resource as below.

```
» cat main.tf
```

```

1 module "todoapp-backend" {
2   source = "git::https://github.com/CSSE6400/terraform//container"

4   image = "ghcr.io/csse6400/todo-app:combined-latest"
5   instance_type = "t2.micro"
6   environment = {
7     APP_ENV="local"
8     APP_KEY="base64:8PQEPYGlTm1t3aqWmlAw/ZPwCiIFvdXDBjk3mhsom/A="
9     APP_DEBUG="true"
10    LOG_LEVEL="debug"
11    DB_CONNECTION="mysql"
12    DB_HOST=aws_db_instance.todoapp-database.address
13    DB_PORT="3306"
14    DB_DATABASE="todoapp"
15    DB_USERNAME="todoapp"
16    DB_PASSWORD=local.password
17  }
18  ports = {
19    "80" = "8000"
20  }
21  security_groups = [aws_security_group.todoapp-backend.name]

23  tags = {
24    Name = "todoapp-backend"
25  }
26 }

```

Note that we are passing the address of our remote database into the container as an environment variable. This is a module which requires a source. In our case, the source will be the Github repository created earlier. Also notice that we map port 80 of the EC2 machine to port 8000 within the container, we should create a security group to make the instance accessible.

```

» cat main.tf

1 resource "aws_security_group" "todoapp-backend" {
2   name = "todoapp-backend"
3   description = "Todo App HTTP and SSH access"

5   ingress {
6     from_port = 80
7     to_port = 80
8     protocol = "tcp"
9     cidr_blocks = ["0.0.0.0/0"]
10  }

12  ingress {
13    from_port = 22
14    to_port = 22
15    protocol = "tcp"

```

```

16     cidr_blocks = ["0.0.0.0/0"]
17 }
18
19 egress {
20     from_port = 0
21     to_port = 0
22     protocol = "-1"
23     cidr_blocks = ["0.0.0.0/0"]
24 }
25 }

```

You will also want to create an output block to expose the address of the instance.

```

» cat main.tf
1 output "url" {
2     value = module.todoapp-backend.public_dns
3 }

```

This should give you a `main.tf` file which fully deploys a todo application. If you haven't been applying as we go, try and apply the Terraform file now. If you have any issues, ask your tutor for guidance.

References

- [1] M. Kleppmann, *Designing Data-Intensive Applications: The big ideas behind reliable, scalable, and maintainable systems*. O'Reilly Media, Inc., March 2017.
- [2] B. Webb, "Infrastructure as code," March 2022. <https://csse6400.uqcloud.net/handouts/iac.pdf>.