

Scaling Stateless Components

Software Architecture

March 29, 2023

Evan Hughes & Anna Truffet & Brae Webb



1 This Week

Our goal is to scale out the stateless component of our TaskOverflow application across multiple compute instances. Specifically we will need to:

- Route traffic to our deployed TaskOverflow application with a *load balancer*.
- Scale out TaskOverflow instances with *autoscaling*.
- Check the status of our instances with a *healthcheck*.
- *Dynamically scale* our application based on load.

2 Load Balancers

Load balancing distributes a load over a set of resources. For example, balancing network traffic across several servers. Load balancing is crucial to the scalability of modern systems, as often, one physical device cannot handle the volume of requests or the processing demand (e.g. the large amount of network traffic for a large website).

A service which load balances, is called a **Load Balancer**.

2.1 Routing Algorithms

A load balancer can implement many techniques to select which resource to route incoming requests toward, these techniques are the load balancer's routing algorithm.

Below are several common routing techniques. There are many other generic and bespoke routing algorithms that are not listed.

Round Robin allocates requests to the next available server regardless of where the last request was sent. It is simple, and in practice, works effectively.

Least Connections sends the next request to the node with the fewest current connections. The load balancer is responsible for tracking how many connections exist to each node.

Weighted Least Connections sends the next request to the node with the least weighted connections. This is similar to the above least connections method, however, each node has an associated weight. This allows certain nodes to be preferred over others. This is useful if we have an unequal distribution of compute power. We would want to give smaller nodes a reduced load in comparison to other more powerful nodes.

Consistent Hashing In some cases we may want a user to consistently be routed to a specific node. This is useful for multiple transactions that need to be done in a consistent order or if the data is stored/-cached on the node. This can be done by hashing the information in the request payload or headers and then routing the request to the node that handles hashes in the range of the computed hash.

2.2 Health Checks

When load balancing, it is important to ensure that the nodes to which we route requests are able to service the request. A good health check can save or break your service. Consider the two following examples from UQ's Information Technology Services (ITS):

Example 1 Early in my career, I, Evan Hughes, setup a multi-node Directory Server at UQ under the hostname of `ldap.uq.edu.au`. This server was a NoSQL database which implemented the LDAP protocol and supported UQ Authenticate, UQ's Single Sign-On service.

The service had a load balancer which checked that port 389 was open and reachable. This worked well most of the time. However, the health check was too weak. When

1. a data-center outage occurred; and
2. the storage running the service disappeared; but
3. the service was still running in memory.

The health check passed, but in reality, the service was talking to dead nodes, causing upstream services to have intermittent failures.

Example 2 During the rollout of a new prompt for UQ Authenticate, which required users to go to my.UQ to provide verified contact details – the Blackboard (learn.uq.edu.au) service went completely offline. The health check for Blackboard at the time completed a full authentication as a test user to ensure everything was functioning as expected. Once this user was enrolled into the new rollout, the health checks started reporting failures and within a matter of minutes the entire pool of nodes were shutdown. This health check was too broad and was not isolated enough to the service that it was checking.

A lot of services will provide a health check endpoint or a metrics endpoint which can help the engineer setup a proper level of health check. We want a health check that is specific enough for the service that it is checking but not so specific that it is too brittle. For the TaskOverflow application that we have been building so far, a reasonable health check would be that the health endpoint ensures the database is available and that the application is able to connect to it.

3 Load Balancers in AWS

3.1 Types of AWS Load Balancers

Not all load balancers are the same. Some load balancers inspect the transmitted packets to correctly route the packet. We will cover two types of load balancers provided by AWS:

Application Load Balancer is an OSI layer 7¹ load balancer which routes traffic based on the request's content. This is useful for services using HTTP, HTTPS, or any other supported protocol.

Network Load Balancer is an OSI layer 4² load balancer which routes traffic based on the source and destination IP addresses and ports. This is useful for services that are using TCP or UDP.

3.2 AWS Load Balancer Design

An AWS Elastic³ Load Balancer has three distinct components.

Listeners allows traffic to enter the Elastic Load Balancer. Each listener has a port (e.g. port 80) and a protocol (e.g. HTTP) associated with it.

Target Groups are groups of nodes which the load balancer can route to. Each target group has a protocol and a port associated with it, allowing us (the programmer) to switch ports on the way through the load balancer. This is useful if the targets are using a different port to the ports we want to expose.

Load Balancer is the actual load balancer that routes the traffic to the target groups based on rules that we setup. The load balancer has a DNS name that we can use to route traffic to it. The load balancer also has a security group that we can use to control what traffic can enter the load balancer.

¹OSI layer 7: Application, in this case HTTP/HTTPS/etc

²OSI layer 4: Transport, in this case TCP/UDP

³Elastic, in a cloud computing context, refers to the system's ability to adapt to workload by starting and stopping infrastructure services to meet demand.

AWS Application Load Balancer Components



3.3 Autoscaling in AWS

Instead of creating the maximum amount of services we predict we will need, we can automatically scale the number of nodes we need to minimise resources. When the load is low, we operate with minimal nodes. When the load is high, we increase the number of nodes available to cope.

To compute the resources needed, AWS relies on triggers from CloudWatch and scaling policies. Some premade triggers are based around a node's

- CPU usage,
- memory usage, or
- network usage.

We create custom triggers based on our application's metrics.

4 Load Balancing TaskOverflow

This week we are going to explore load balancing the TaskOverflow service that we have been working with. The aim is to have a service that, when given a lot of requests, will be able to scale out the web server instances to handle it. This will not be a full solution to the scaling issue as our database is still a single node, but it will be a good start.⁴ Other methods could also be employed to help deal with the load like caching but we will leave that for another day.

⁴We will not explore scaling persistent data in the practicals. If you wish to try in your assignment, please see some of the concepts presented in the Distributed Systems II lecture [1].

Getting Started

1. Using the GitHub Classroom link for this practical provided by your tutor in Slack, create a repository to work within.
2. Install Terraform if not already installed, as it will be required again this week.
3. Start your learner lab and copy the AWS Learner Lab credentials into a credentials file in the root of the repository.
4. In the repository you will have a given folder with two subfolders. `ec2` is the starter Terraform code for the EC2 path from last week. `ecs` is the starter Terraform code for the ECS path from last week.
5. Copy the contents of the folder into the root of the repository based on the path you wish to take.

As per last week, we encourage that you take Path B first and then peruse Path A, if you have time. This week, Path A offers an alternative method to autoscaling.

4.1 [Path A] EC2

Congratulations! You have chosen to go down the EC2 path. This week we need to create an AutoScaling Group and a Load Balancer to handle the scaling of our service. Our goal is to implement the deployment diagram below.



To get started, we need to modify how we create our EC2 instances. Instead of manually deploying an instance we are going to tell AWS how to deploy instances for us. We can do this by converting our EC2 instance into a Launch Template. Our launch template is going to look very similar to the instance we had last week, with a few variables changing. This includes that `user_data` must be stored as a base64 encoded string. Some of these changes have been highlighted below.

```

resource "aws_launch_template" "todo" {
  name          = "todo-launch-template"
  image_id      = "ami-005f9685cb30f234b"
  instance_type = "t2.micro"
  key_name      = "vockey"
  user_data     = base64encode(<<-EOT
    #!/bin/bash
    yum update -y
    yum install -y docker
    service docker start
    systemctl enable docker
    usermod -a -G docker ec2-user
    docker run --restart always -e SQLALCHEMY_DATABASE_URI=postgresql://${local.
    database_username}:${local.database_password}@${aws_db_instance.database.address
    }:5432/todo -p 6400:6400 ${local.image}
EOT
  )
}

vpc_security_group_ids = [aws_security_group.todo.id]
}

resource "aws_security_group" "todo" {
  name          = "todo"
  description   = "TaskOverflow Security Group"

  ingress {
    from_port   = 6400
    to_port     = 6400
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

```

Now that we have an instance template and a security group, we can create the AutoScaling Group.

The AutoScaling group is given the instance template and uses it to create groups of instances. The group defines the desired capacity and the limits of how many instances we want.

Building a AWS Load Balancer



```
» cat autoscaling.tf

resource "aws_autoscaling_group" "todo" {
  name          = "todo"
  availability_zones = ["us-east-1a"]
  desired_capacity    = 1
  max_size          = 4
  min_size          = 1

  launch_template {
    id      = aws_launch_template.todo.id
    version = "$Latest"
  }
}
```

If we deployed this now, we would get one initial instance created for us, based on the `desired_capacity` and `min_size` values. Now that we have instances we have the ability for traffic to be routed to them. Let's create a target group which will be on the internal facing side of our load balancer. This is the side that lets traffic from the load balancer hit the instances.

Building a AWS Load Balancer



```
» cat lb.tf

resource "aws_lb_target_group" "todo" {
  name      = "todo"
  port      = 6400
  protocol = "HTTP"
  vpc_id    = aws_security_group.todo.vpc_id
}
```

We also need to link the target_group to the autoscaling group.

```
» cat lb.tf

resource "aws_autoscaling_attachment" "todo" {
  autoscaling_group_name = aws_autoscaling_group.todo.id
  alb_target_group_arn   = aws_lb_target_group.todo.arn
}
```

We now need to create a load balancer and route traffic to our target group. You will notice that there is also a security group attached to the load balancer. This is the firewall that sits on the front of the load balancer and allows traffic to enter the load balancer.

Building a AWS Load Balancer



```
» cat lb.tf
```

```
resource "aws_lb" "taskoverflow" {
  name          = "taskoverflow"
  internal      = false
  load_balancer_type = "application"
  subnets       = data.aws_subnets.private.ids
  security_groups = [aws_security_group.taskoverflow.id]
}

resource "aws_security_group" "taskoverflow" {
  name          = "taskoverflow"
  description   = "TaskOverflow Security Group"

  ingress {
    from_port    = 80
    to_port      = 80
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

The load balancer now needs a way to accept incoming traffic. We need a listener which creates the rules for how traffic is routed to the instances.

Building a AWS Load Balancer



```
`` cat lb.tf ``

resource "aws_lb_listener" "taskoverflow" {
  load_balancer_arn = aws_lb.taskoverflow.arn
  port             = "80"
  protocol         = "HTTP"

  default_action {
    type           = "forward"
    target_group_arn = aws_lb_target_group.todo.arn
  }
}
```

Now we are able to deploy our infrastructure.

```
$ terraform apply
```

We have created the architecture as requested with a dynamically scaling group of EC2s and an Application Load Balancer in front which is listens on port 80 and routes to port 6400 of the instances.

To make our autoscaling group perform dynamic scaling, we need to create some rules for when to scale up and when to scale down. The policies below track the average CPU usage of the group of instances. If the usage is above 20 percent it will scale up by 1 instance, and if it is below 10 percent it will scale down by 1 instance. We also apply a cooldown period so that we do not act too aggressively to changes in CPU usage, as instances will take time to start and handle load.

```

resource "aws_autoscaling_policy" "todo_scale_down" {
  name          = "todo_scale_down"
  autoscaling_group_name = aws_autoscaling_group.todo.name
  adjustment_type      = "ChangeInCapacity"
  scaling_adjustment    = -1
  cooldown           = 120
}

resource "aws_cloudwatch_metric_alarm" "todo_scale_down" {
  alarm_description  = "Monitors CPU utilization for Todo"
  alarm_actions      = [aws_autoscaling_policy.todo_scale_down.arn]
  alarm_name         = "todo_scale_down"
  comparison_operator = "LessThanOrEqualToThreshold"
  namespace          = "AWS/EC2"
  metric_name        = "CPUUtilization"
  threshold          = "10"
  evaluation_periods = "2"
  period              = "120"
  statistic           = "Average"

  dimensions          = {
    AutoScalingGroupName = aws_autoscaling_group.todo.name
  }
}

resource "aws_autoscaling_policy" "todo_scale_up" {
  name          = "todo_scale_up"
  autoscaling_group_name = aws_autoscaling_group.todo.name
  adjustment_type      = "ChangeInCapacity"
  scaling_adjustment    = 1
  cooldown           = 120
}

resource "aws_cloudwatch_metric_alarm" "todo_scale_up" {
  alarm_description  = "Monitors CPU utilization for Todo"
  alarm_actions      = [aws_autoscaling_policy.todo_scale_up.arn]
  alarm_name         = "todo_scale_up"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  namespace          = "AWS/EC2"
  metric_name        = "CPUUtilization"
  threshold          = "20"
  evaluation_periods = "2"
  period              = "120"
  statistic           = "Average"

  dimensions          = {
    AutoScalingGroupName = aws_autoscaling_group.todo.name
  }
}

```

Head to Section 4.3 to see how we can send multiple requests to our service so that the load balancer has to scale up.

4.2 [Path B] ECS

Congratulations! You have chosen to go down the ECS path. This week we need to create an Application Scaling Policy for our ECS service and a Load Balancer to handle the routing of our service. Our goal is to implement the deployment diagram below.



Last week, when we setup the ECS service, we noticed that we could not get an endpoint because the instance would only be provisioned after our Terraform code had run. This is because ECS is already a dynamic scaling service and it expects a load balancer to route its traffic. To get started we need to create a target group which defines where traffic can be routed to.

Building a AWS Load Balancer



```
» cat lb.tf

resource "aws_lb_target_group" "todo" {
  name        = "todo"
  port        = 6400
  protocol    = "HTTP"
  vpc_id      = aws_security_group.todo.vpc_id
  target_type = "ip"

  health_check {
    path          = "/api/v1/health"
    port          = "6400"
    protocol     = "HTTP"
    healthy_threshold = 2
    unhealthy_threshold = 2
    timeout      = 5
    interval     = 10
  }
}
```

Load balancing is core to how ECS works and so the `aws_ecs_service` resource that we used last week accepts a `load_balancer` block. To associate the target group with our ECS service, modify the given `aws_ecs_service.taskoverflow` resource to include a `load_balancer` block.

```
» cat ecs.tf

load_balancer {
  target_group_arn = aws_lb_target_group.todo.arn
  container_name   = "todo"
  container_port   = 6400
}
```

With the internal side of the load balancer done, we can create it and a firewall for the external side. This firewall allows us to restrict what traffic will be allowed to reach the load balancer.

Building a AWS Load Balancer



```
» cat lb.tf
```

```
resource "aws_lb" "taskoverflow" {
  name          = "taskoverflow"
  internal      = false
  load_balancer_type = "application"
  subnets       = data.aws_subnets.private.ids
  security_groups = [aws_security_group.taskoverflow.id]
}

resource "aws_security_group" "taskoverflow" {
  name          = "taskoverflow"
  description   = "TaskOverflow Security Group"

  ingress {
    from_port    = 80
    to_port      = 80
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

Now over to the external side of the load balancer. We need to create a listener which is the entry point for the load balancer.

Building a AWS Load Balancer



```

» cat lb.tf

resource "aws_lb_listener" "todo" {
  load_balancer_arn    = aws_lb.taskoverflow.arn
  port                  = "80"
  protocol              = "HTTP"

  default_action {
    type      = "forward"
    target_group_arn = aws_lb_target_group.todo.arn
  }
}
  
```

If we deployed now, we would have implemented the deployment diagram above. However, we want to add autoscaling to our service so that it can scale up and down based on the load.

```

» cat autoscaling.tf

resource "aws_appautoscaling_target" "todo" {
  max_capacity      = 4
  min_capacity      = 1
  resource_id        = "service/taskoverflow/taskoverflow"
  
```

```

scalable_dimension  = "ecs:service:DesiredCount"
service_namespace   = "ecs"
}

resource "aws_appautoscaling_policy" "todo-cpu" {
  name          = "todo-cpu"
  policy_type   = "TargetTrackingScaling"
  resource_id    = aws_appautoscaling_target.todo.resource_id
  scalable_dimension = aws_appautoscaling_target.todo.scalable_dimension
  service_namespace = aws_appautoscaling_target.todo.service_namespace

  target_tracking_scaling_policy_configuration {
    predefined_metric_specification {
      predefined_metric_type = "ECSServiceAverageCPUUtilization"
    }

    target_value           = 20
  }
}

```

This auto scaling policy looks at the average CPU utilization of the service and scales up if it is above 20% and scales down if it is below 20%. This is a very simple policy but it is a good starting point. We can now deploy our service and see it scale up and down.

Continue onto the next section to see how to send multiple requests to our service to generate traffic to trigger scaling.

4.3 Producing Load with k6

We have a service but us visiting it in our web browser is not going to be enough load for our scaling policies to trigger. To do this we will employ the help of a tool called k6, which is a relatively new load testing tool. To install k6 visit <https://k6.io/docs/get-started/installation/>. It can be installed in the code spaces environment or locally.

We have provided an example k6 file, which is JavaScript code that creates 1000 to 5000 users to call the list endpoint of our service.

```

» cat k6.js

import http from 'k6/http';
import { sleep, check } from 'k6';

export const options = {
  stages: [
    { target: 1000, duration: '1m' },
    { target: 5000, duration: '10m' },
  ],
};

export default function () {

```

```
const res = http.get('http://your-loadBalancer-url-here/api/v1/todos');
check(res, { 'status was 200': (r) => r.status == 200 });
sleep(1);
}
```

We can then run this file using the following command.

```
$ k6 run k6.js
```

```
execution: local
script: load.js
output: -

scenarios: (100.00%) 1 scenario, 5000 max VUs, 11m30s max duration (incl. graceful
stop):
  * default: Up to 5000 looping VUs for 11m0s over 2 stages (gracefulRampDown:
  30s, gracefulStop: 30s)

running (00m05.4s), 0091/5000 VUs, 140 complete and 0 interrupted iterations
default  [-----] 0091/5000 VUs 00m05.4s/11m00.0s
```

4.3.1 EC2 Auto Scaling

With all the pieces together we can now see if our efforts have paid off. While the above k6 code is running, let's check the EC2 console and see what actions our autoscaling policy has taken. In the EC2 console scroll down the left hand menu and select Auto Scaling Groups.

The screenshot shows the AWS EC2 Home page. The left sidebar contains a navigation menu with the following items:

- Dedicated Hosts
- Scheduled Instances
- Capacity Reservations
- Images** (selected)
- AMIs
- AMI Catalog
- Elastic Block Store**
- Volumes
- Snapshots
- Lifecycle Manager
- Network & Security**
- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs
- Network Interfaces
- Load Balancing**
- Load Balancers
- Target Groups
- Auto Scaling** (selected)
- Launch Configurations
- Auto Scaling Groups

The main content area displays a summary of Amazon EC2 resources in the US:

Instances (running)	0	Auto Scaling Group
Elastic IPs	0	Instances
Load balancers	1	Placement groups
Snapshots	0	Volumes

A callout box highlights the "Auto Scaling Group" link.

Launch instance
To get started, launch an Amazon EC2 instance, which is a virtual server in the cloud.

Launch instance (button)
Migrate a server (button)

Note: Your instances will launch in the US East (N. Virginia) Region

Scheduled events

You will be presented with our todo group, which states the current instances and the desired minimum and maximum number of instances. Select the name of the group.

The screenshot shows the AWS Auto Scaling groups page. The top navigation bar includes links for EC2 and Auto Scaling groups.

Auto Scaling groups (1) [Info](#)

Create an Auto Scaling group

Search your Auto Scaling groups

todo | Version Latest

Name	Launch template/configuration	Instances	Status	Desired capacity	Min
todo	todo-launch-template	1	-	1	1

In this panel we are presented with much of the same information, but what we want is the Automatic scaling tab.

todo

Details Activity Automatic scaling Instance management Monitoring Instance refresh

Group details

Edit

Auto Scaling group name todo	Desired capacity 1	Status -	Amazon Resource Name (ARN) arn:aws:autoscaling:us-east-1:24517 7958106:autoScalingGroup:91e8402c-a8 be-4c4f-ad6f-0c02b19607fc:autoScaling GroupName/todo
Date created Mon Mar 27 2023 16:32:28 GMT+1000 (Australian Eastern Standard Time)	Minimum capacity 1	Maximum capacity 4	

In this tab our two policies are displayed and we can see the configuration of both.

todo

Details Activity **Automatic scaling** Instance management Monitoring Instance refresh

Scaling policies resize your Auto Scaling group to meet changes in demand. With reactive dynamic scaling policies, you can track specific CloudWatch metrics and take action when the CloudWatch alarm threshold is met. Use predictive scaling policies along with dynamic scaling policies in the following situations: when your application demand changes quickly, but with a recurring pattern, or when your EC2 instances require more time to initialize.

Dynamic scaling policies (2) [Info](#)



Actions ▾

Create dynamic scaling policy

< 1 >

todo_scale_down



Simple scaling

Enabled

todo_scale_down

breaches the alarm threshold: CPUUtilization <= 10 for 2 consecutive periods of 120 seconds for the metric dimensions:

AutoScalingGroupName = todo

Remove 1 capacity units

120 seconds before allowing another scaling activity

todo_scale_up



Simple scaling

Enabled

todo_scale_up

breaches the alarm threshold: CPUUtilization > 20 for 2 consecutive periods of 120 seconds for the metric dimensions:

AutoScalingGroupName = todo

Add 1 capacity units

120 seconds before allowing another scaling activity

In the Activity tab we can see the actions performed by our policies, whether it be a scale up or a scale down.

todo

Details	Activity	Automatic scaling	Instance management	Monitoring	Instance refresh
Activity notifications (0)					
<input type="text"/> Filter notifications	<input type="button"/>	<input type="button"/> Actions ▾	<input type="button"/> Create notification	< 1 >	<input type="button"/>
Send to On instance action					
Activity history (2)					
<input type="text"/> Filter activity history	<input type="button"/>	< 1 >	<input type="button"/>	< 1 >	<input type="button"/>
Status	Description	Cause			Start time
Successful	Updating load balancers/target groups: Successful. Status Reason: Added: arn:aws:elasticloadbalancing:us-east-1:245177958106:targetgroup/todo/ebcbca97edb519ec (Target Group).				2023 March 27, 04:33:20 PM +10:00
Successful	Launching a new EC2 instance: i-05bf864d60fa54dbe	At 2023-03-27T06:32:28Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 1. At 2023-03-27T06:32:31Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.			2023 March 27, 04:32:33 PM +10:00

If we head over to CloudWatch alarms, we can see the alarms that are the triggers for our policies.



4.3.2 ECS Auto Scaling

With all the pieces together we can now see if our efforts have paid off. While the k6 code from above is running, let's go to the ECS console and see if we can see any scaling events. Navigate to ECS -> Clusters -> taskoverflow -> Services -> taskoverflow -> Configuration and tasks.

taskoverflow Info



Health and metrics

Logs

Configuration and tasks

Deployments and events

Networking

Tags

In this panel we can see our Auto Scaling configuration which lists the desired, minimum and maximum number of tasks. The policies are listed further down and we can view the alarm which is a graphical representation of the CPU utilization vs the target value.

Auto Scaling

Desired tasks 1	Min tasks 1	Max tasks 4
--------------------	----------------	----------------

Policies (1)

Policy name	Policy type	Scale-in	Alarm
todo-cpu: Tracking ECSServiceAverageCPUUtilization at 20	Target tracking	On	TargetTracking-service/taskoverflow /taskoverflow-AlarmHigh-a56c53cb-c376-4ff6-a18a-64adaec37997

Tasks (1/2)

Task	Last status	Desired st...	Tas...	Revision	Health sta...	Started at	Container insta...
● 132a0...	⌚ Running	⌚ Running	todo	6	 ⓘ Unknown	40 seconds ago	-

In the Cloudwatch Alarm panel you will notice that we have two different alarms, these are for the scaling up and down of the service. Selecting the alarm you can view the status where an "in alert" alarm is when the auto scaling configuration needs to action increasing/decreasing the number of instances.



5 Conclusion

You have now deployed an scalable stateless service. If you would like to experiment with generating different load for the service, please read through the k6 documentation at <https://k6.io/docs/>.

In the cloud assignment, we will use k6 to test various scenarios as described in the task sheet and evaluate how your service performs. It will be beneficial to be familiar with how this type of testing works.

References

- [1] B. Webb, “Distributed systems II slides,” March 2023. <https://csse6400.uqcloud.net/slides/distributed2.pdf>.