

Distributed Systems I

Software Architecture

Brae Webb & Richard Thomas

March 23, 2026



Mathias Verras
@mathiasverraes

There are only two hard problems
in distributed systems:

2. Exactly-once delivery
1. Guaranteed order of messages
2. Exactly-once delivery

Lecture Goal: Balance a healthy love-hate relationship with dis-
tributed systems

Going forward

Investigating architectures that are *distributed*.

Distributed Systems Series

Distributed I *Reliability* and *scalability* of
stateless systems.

Distributed II *Complexities* of *stateful*
systems.

Distributed III *Hard problems* in distributed
systems.

What are the benefits?

- Improved *reliability*
 - Improved *scalability*
 - Improved *latency*
- Some systems are inherently distributed.

What are the drawbacks?

- Increased *complexity*
- Increased *attack surface*
- Increased *latency*
- Introduce *consistency* problems

We'll look at a few reasons that distributed systems are *fundamentally* quite challenging

§ Fallacies

A few reasons for complexity

Fallacies of *Distributed Computing*

Sun Microsystems in 1994, primarily accredited to Peter Deutsch
(doy-ch)

Fallacy #1

The network is reliable



Success: Send request to add item to cart.



Failure 1: Request not received by server (CartService).



Failure 1: *Solution* resend request?



If the service goes down and all clients are re-trying,
the service is in for a shock when it comes back,
we solve this with *exponential backoff*.

Exponential Backoff

```
1  retries = 0
2  do:
3      status = service.request()
5
5      if status != SUCCESS:
6          retries += 1
7          wait(2 ** retries)
8  while (status != SUCCESS and retries < MAX_RETRIES)
```



Failure 2: Response not received.



- Failure 1's *solution* of resending request leads to:
- *Duplicate actions*, problem for ordering/payments.



Fallacy #2

Latency is zero

Network Statistics

Home to UQ

Home to us-east-1

EC2 to EC2

Network Statistics

Home to UQ 20.025ms

Home to us-east-1

EC2 to EC2

Network Statistics

Home to UQ 20.025ms

Home to us-east-1 249.296ms

EC2 to EC2

Network Statistics

Home to UQ 20.025ms

Home to us-east-1 249.296ms

EC2 to EC2 0.662ms

- Be mindful when designing distributed systems.
- Network calls are *much* slower than local calls.

Fallacy #3

Bandwidth is infinite

Similar to previous fallacy, be mindful,
distributed calls clog up network.

Definition 0. Stamp Coupling

Components which share a composite data structure.

Particularly an issue if components are on different compute nodes.

Fallacy #4

The network is secure



Authentication only occurs when entering Sahara data centre.



- Bad actor gets access via one insecure node.
- Network is *compromised*.
- Practice *defence in depth*.

Fallacy #5

The topology never changes

- Topology changes all the time, cloud makes this *easier*.
- Don't rely on static IPs.
- Don't assume consistent latency.

Fallacy #6

There is only one administrator

- Things spontaneously break.
- Who can help you?

Fallacy #7

Transport cost is zero

Remember

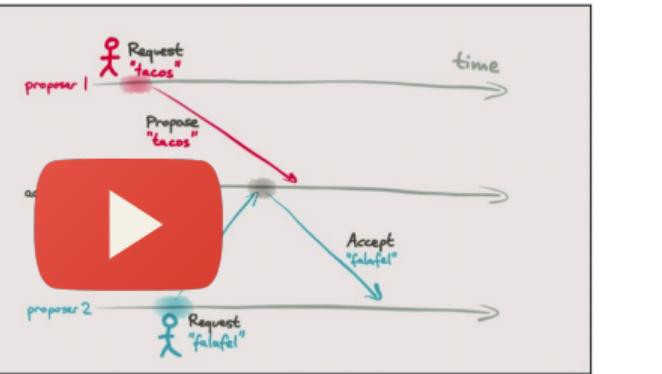
Distributed systems are *hard*. Choosing to use them should be *well considered*.

Can often introduce more problems than they solve.

When you need to, maybe prove it?



Sept 13-14, 2019
thestrangeloop.com



- Presentation: Proving correctness of distributed systems.
- by Dr. Martin Kleppmann, University of Cambridge.
- Using Isabelle – See CSS research sub-group.

Or, more realistically,

Use *existing* algorithms and software

Distributed Systems Series

Distributed I *Reliability* and *scalability* of
stateless systems.

Distributed II Complexities of stateful
systems.

Distributed III Hard problems in distributed
systems.

Stateless vs. Stateful Systems

Stateless Does *not* utilise *persistent data*.

Stateful Does utilise *persistent data*.

Question

What makes software *reliable*?

Definition 0. Reliable Software

Continues to work, even when things go wrong.

Definition 0. Fault

Something goes wrong.

Death, taxes, and computer system failure are all inevitable to some degree.

Plan for the event.

– Howard and LeBlanc

*Reliable software is
Fault **tolerant***

John von Neumann built fault tolerant hardware in the 1950s.

Problem

Individual computers fail *all the time*

- 10-50 years hard-drive lifetime.
- 10,000 disks will fail daily.
- Google last had 2.5 million servers.
- ...

Solution

Spread the risk of faults over *multiple computers* or *nodes*

Spreading Risk

If you have software that works with *just one* computer,
spreading the software over *two* computers *halves* the
risk that your software will fail.

Spreading Risk

If you have software that works with *just one* computer,
spreading the software over *two* computers *halves* the
risk that your software will fail.

Adding *100* computers reduces the risk by *100*.



- Why is this software somewhat reliable?
- Any individual service can go down and the rest still work.
- Can we do better?
- Can a service go down but still have that service available?

Question

Who has used *auto-scaling*?

Auto-Scaling Terminology

Auto-scaling group A *collection of instances* managed by auto-scaling.

Auto-Scaling Terminology

Auto-scaling group A *collection of instances* managed by auto-scaling.

Capacity Amount of instances *currently* in an auto-scaling group.

Auto-Scaling Terminology

Auto-scaling group A *collection of instances* managed by auto-scaling.

Capacity Amount of instances *currently* in an auto-scaling group.

Desired Capacity Amount of instances *we want to have* in an auto-scaling group.

Auto-Scaling Terminology

Auto-scaling group A *collection of instances* managed by auto-scaling.

Capacity Amount of instances *currently* in an auto-scaling group.

Desired Capacity Amount of instances *we want to have* in an auto-scaling group.

Scaling Policy How to determine the desired capacity.

Auto-Scaling Terminology

Auto-scaling group A *collection of instances* managed by auto-scaling.

Capacity Amount of instances *currently* in an auto-scaling group.

Desired Capacity Amount of instances *we want to have* in an auto-scaling group.

Scaling Policy How to determine the desired capacity.

Minimum/Maximum Capacity *Hard limits* on the minimum and maximum number of instances.

What we really want

Desired Capacity Amount of *healthy* instances
we want to have in an auto-scaling group.

Health check

Mechanism to determine whether an instance is *healthy*.

Auto-scaling

An example



Product Browsing service keeps going down



We might expect product browsing to have a much higher load than other services



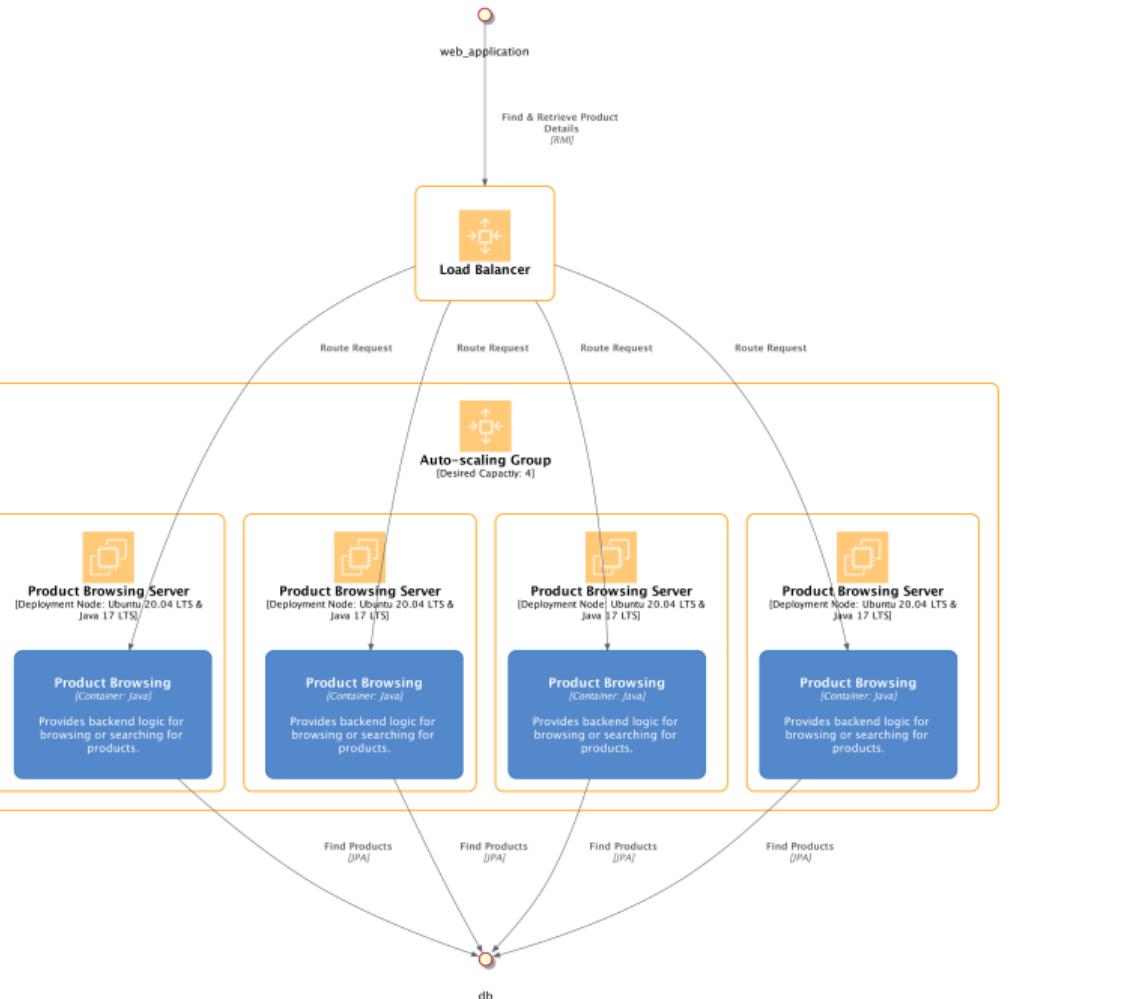
Deploy product browsing on an EC2 instance



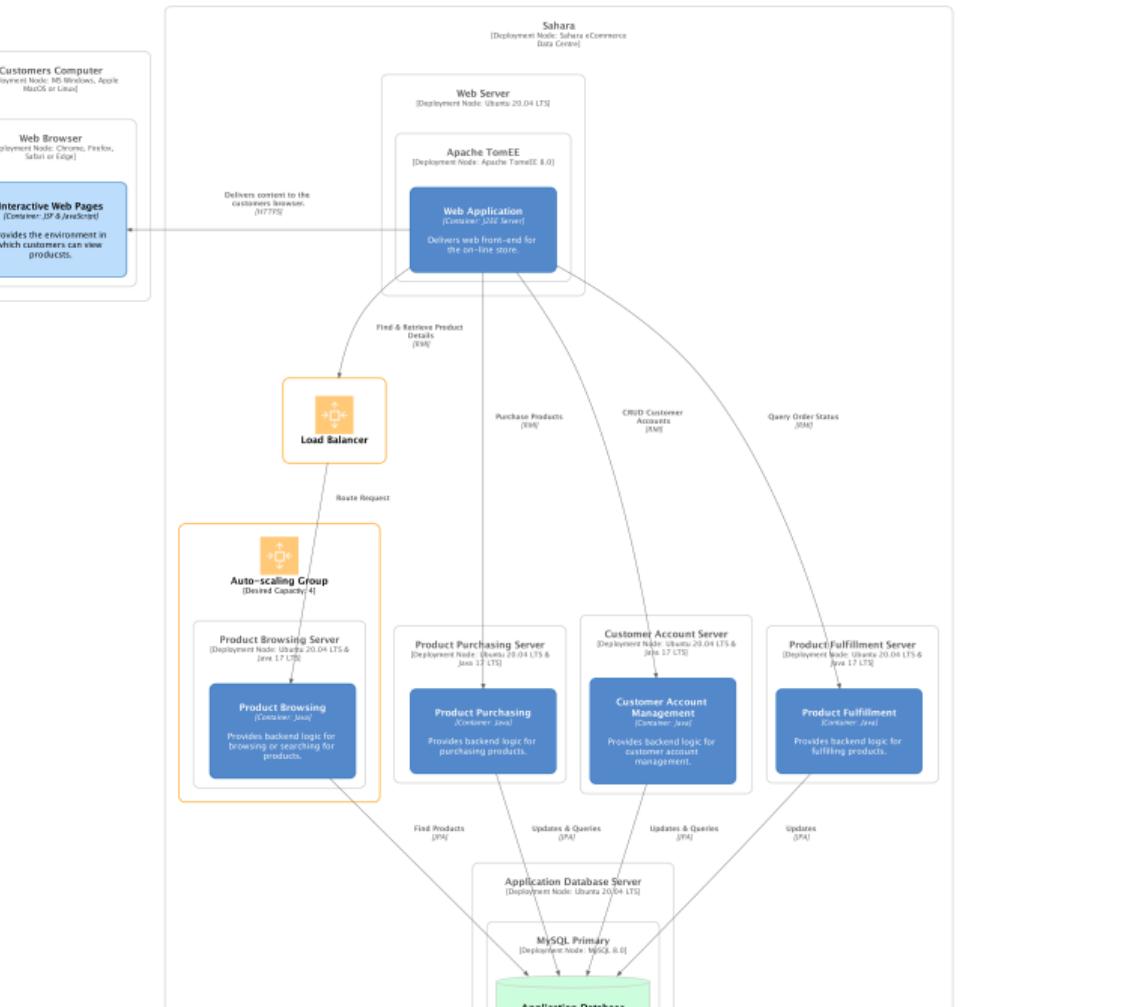
Use an auto-scaling group to replicate the service



What's the problem?



- Traffic was all sent through *one* instance
- Load balancer routes to *all* instances



Load balancing & scaling integrated into overall Sahara deployment diagram

In Summary

Simplicity

Reliability

Scalability

In Summary

Simplicity *Minimal network communication* (compared to other distributed systems), less impacted by fallacies.

Reliability

Scalability

In Summary

Simplicity *Minimal network communication* (compared to *other* distributed systems), less impacted by fallacies.

Reliability Traffic is spread to various services, still *partially operational* if one goes down. Auto-scaling allows for *basic replication*.

Scalability

In Summary

Simplicity *Minimal network communication* (compared to other distributed systems), less impacted by fallacies.

Reliability Traffic is spread to various services, still *partially operational* if one goes down. Auto-scaling allows for *basic replication*.

Scalability Auto-scaling and load balancing allows *individual services to scale*. However, the *database is a bottleneck*.

Database is a bottleneck is foreshadowing