

# Distributed Systems II

*Software Architecture*

Brae Webb & Richard Thomas & Guangdong Bai

March 30, 2026

## *Distributed Systems Series*

Distributed I    *Reliability* and *scalability* of  
*stateless* systems

Distributed II    *Complexities* of *stateful*  
systems

Distributed III    *Hard problems* in distributed  
systems

## *Distributed Systems Series*

Distributed I   Reliability and scalability of stateless systems

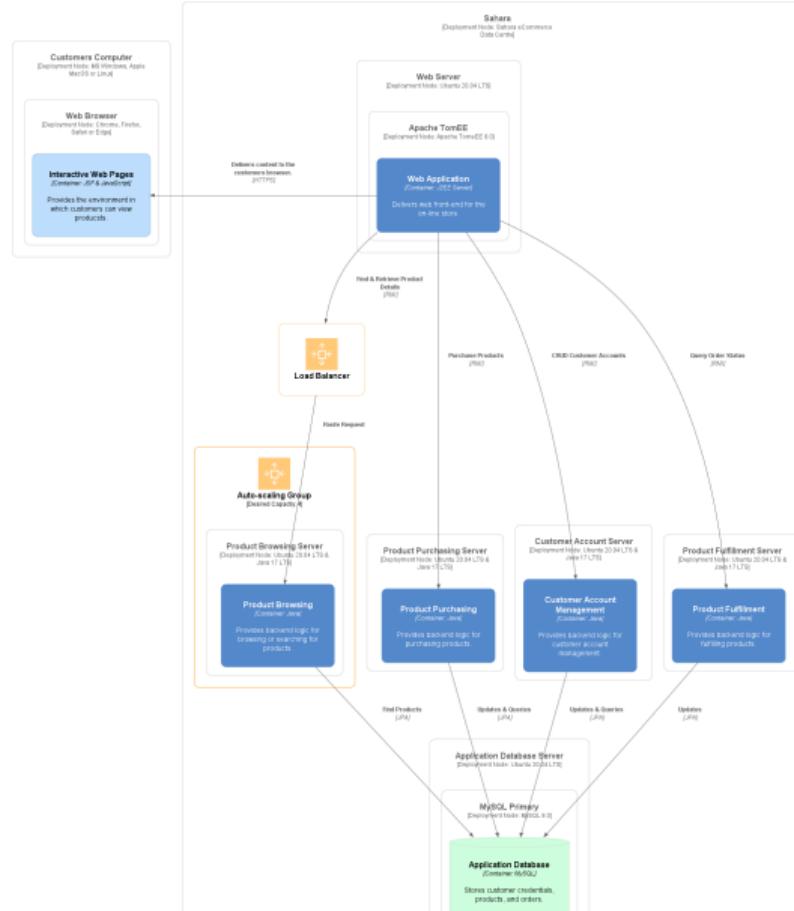
Distributed II   *Complexities* of *stateful* systems

Distributed III Hard problems in distributed systems

*Previously in Distributed I: Benefits...*

- Improved *reliability*
- Improved *scalability*
- Improved *latency*

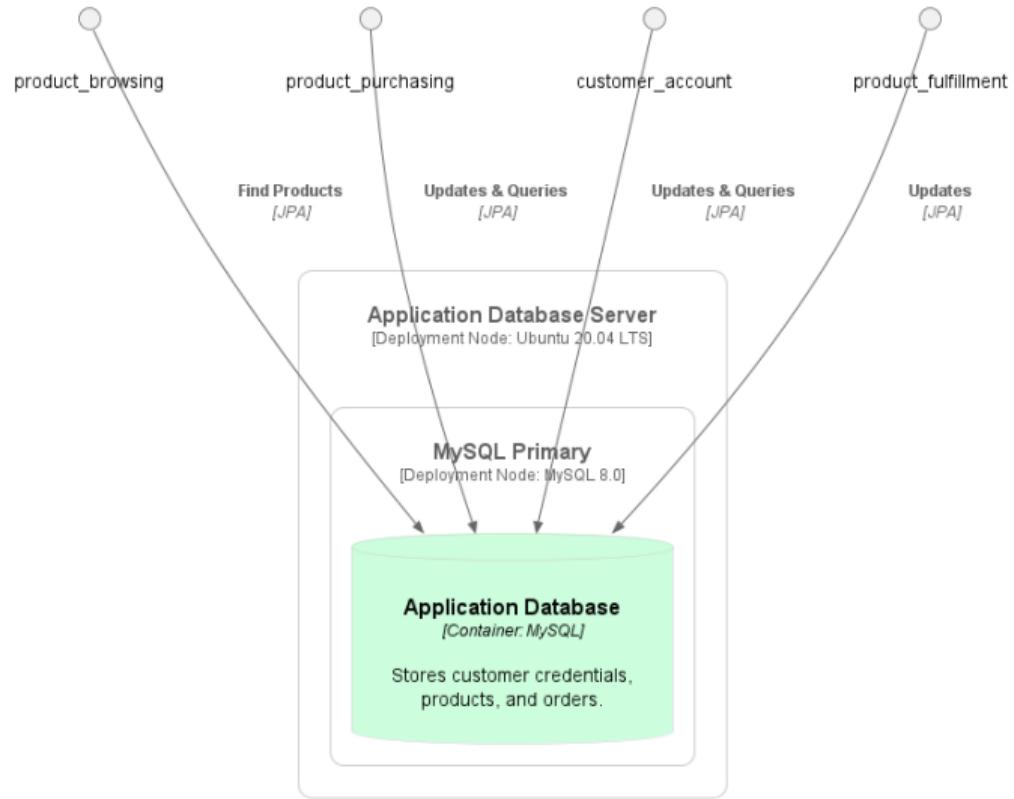
# Previously in Distributed I...



*Question*

What is the *problem*?

# Database



## *Stateless vs. Stateful Systems*

Stateless Does *not* utilise *persistent data*. Or:  
each request is independent.

Stateful Does utilise *persistent data*. Or: the  
server or service remembers and uses  
data from previous interactions.

*Disclaimer*

This is *not* a database course



# Advanced Database Systems (INFS3200)

**Course level**

Undergraduate

**Faculty**

Engineering, Architecture &amp; Information Technology

**School**

Info Tech &amp; Elec Engineering

**Units**

2

**Duration**

One Semester

**Class contact**

2 Lecture hours, 1 Tutorial hour, 1 Practical or Laboratory hour

**Incompatible**

INFS7907

**Prerequisite**

INFS2200

**Assessment methods**

## Current course offerings

Course offerings	Location	Mode	Course Profile
Semester 1, 2022	St Lucia	Internal	<a href="#">COURSE PROFILE</a>
Semester 1, 2022	External	External	<a href="#">COURSE PROFILE</a>
Semester 2, 2022	External	External	PROFILE UNAVAILABLE
Semester 2, 2022	St Lucia	Internal	PROFILE UNAVAILABLE

Please Note: Course profiles marked as not available may still be in development.

## Course description

Distributed database design, query and transaction processing, data integration, data warehousing, data cleansing, management of spatial data, and data from large scale distributed devices.

## Archived offerings

Course offerings	Location	Mode	Course Profile
Semester 1, 2021	St Lucia	Flexible Delivery	<a href="#">COURSE PROFILE</a>
Semester 1, 2021	External	External	<a href="#">COURSE PROFILE</a>
Semester 2, 2021	External	External	<a href="#">COURSE PROFILE</a>
Semester 2, 2021	St Lucia	Internal	<a href="#">COURSE PROFILE</a>
Semester 1, 2020	St Lucia	Internal	<a href="#">COURSE PROFILE</a>

*Question*

How do we fix database scaling issues?

*Question*

How do we fix database scaling issues?

*Answer*

- Replication

*Question*

How do we fix database scaling issues?

*Answer*

- Replication
- Partitioning

*Question*

How do we fix database scaling issues?

*Answer*

- Replication
- Partitioning
- Independent databases

*Question*

How do we fix database scaling issues?

*Answer*

- *Replication*
- Partitioning
- Independent databases

*Question*

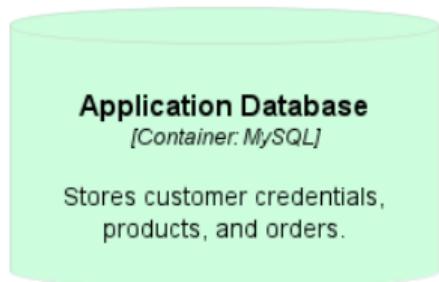
What is *replication*?

*Definition 0.* Replication

Data copied across multiple different machines.



product_id	name	stock	price
1234	Nicholas Cage Reversible Pillow	10	\$10.00
4321	Lifelike Elephant Inflatable	5	\$50.00



product_id	name	stock	price
1234	Nicholas Cage Reversible Pillow	10	\$10.00
4321	Lifelike Elephant Inflatable	5	\$50.00

*Definition 0.* Replica

Database node which stores a copy of the data.

*Question*

What are the advantages of *replication*?

*Question*

What are the advantages of *replication*?

*Answer*

- *Scale* our database to cope with higher loads.

*Question*

What are the advantages of *replication*?

*Answer*

- *Scale* our database to cope with higher loads.
- Provide *fault tolerance* from a single instance failure.

*Question*

What are the advantages of *replication*?

*Answer*

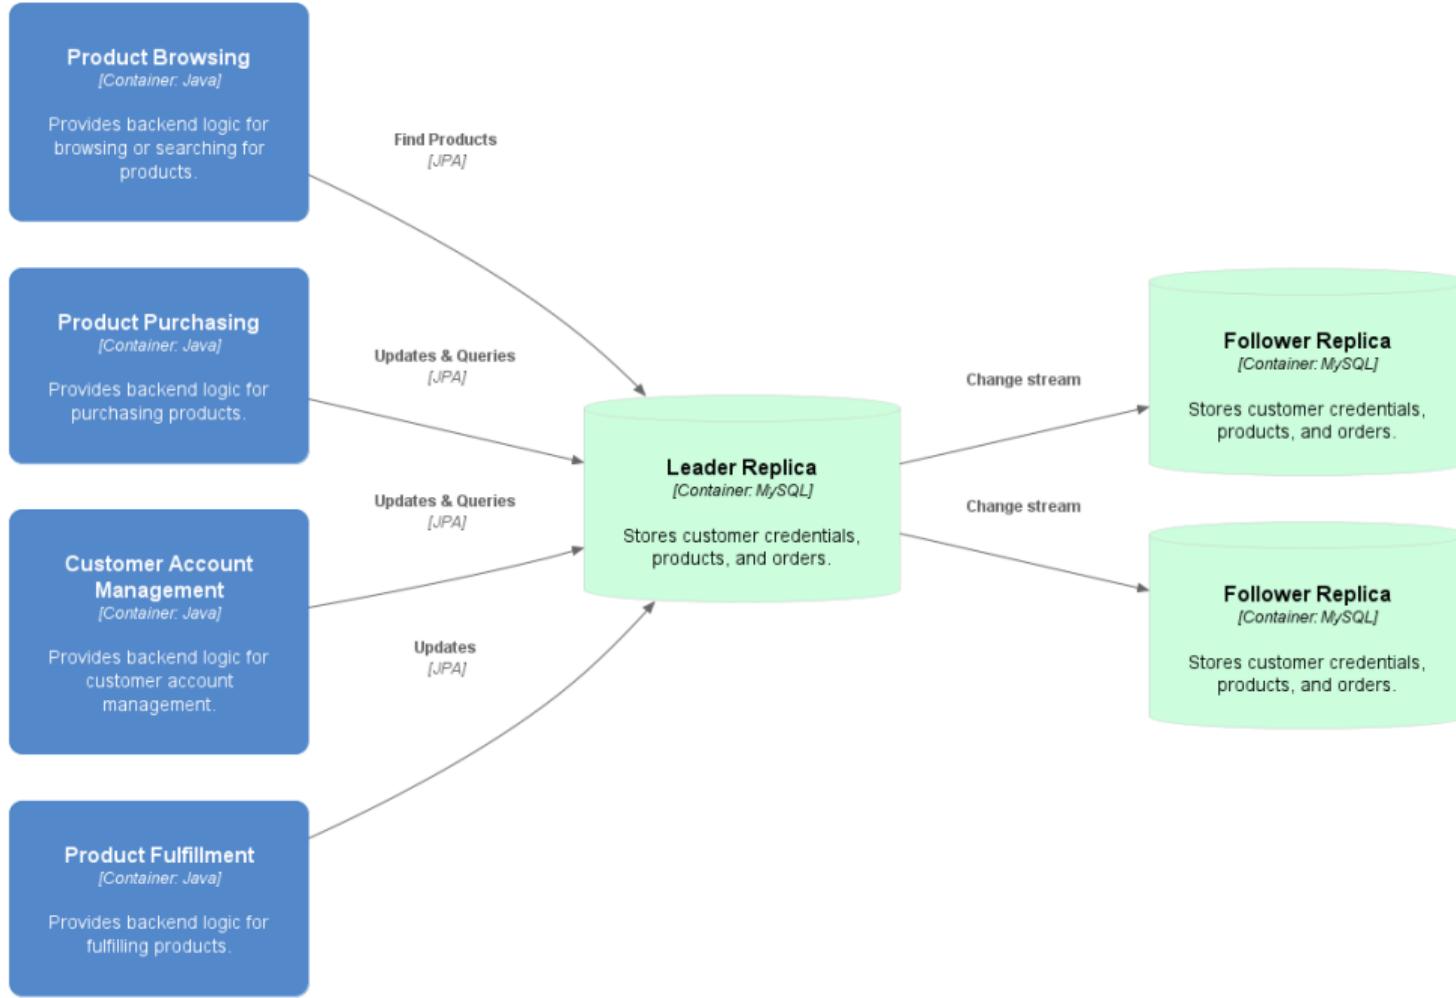
- *Scale* our database to cope with higher loads.
- Provide *fault tolerance* from a single instance failure.
- Locate instances *closer to end-users*.

*Question*

How do we replicate our data?

*First Approach*

Leader-Follower Replication



*Definition 0.* Leader-based Replication

one node (the leader) handles all write operations, and multiple other nodes (followers) replicate the data and handle read operations.

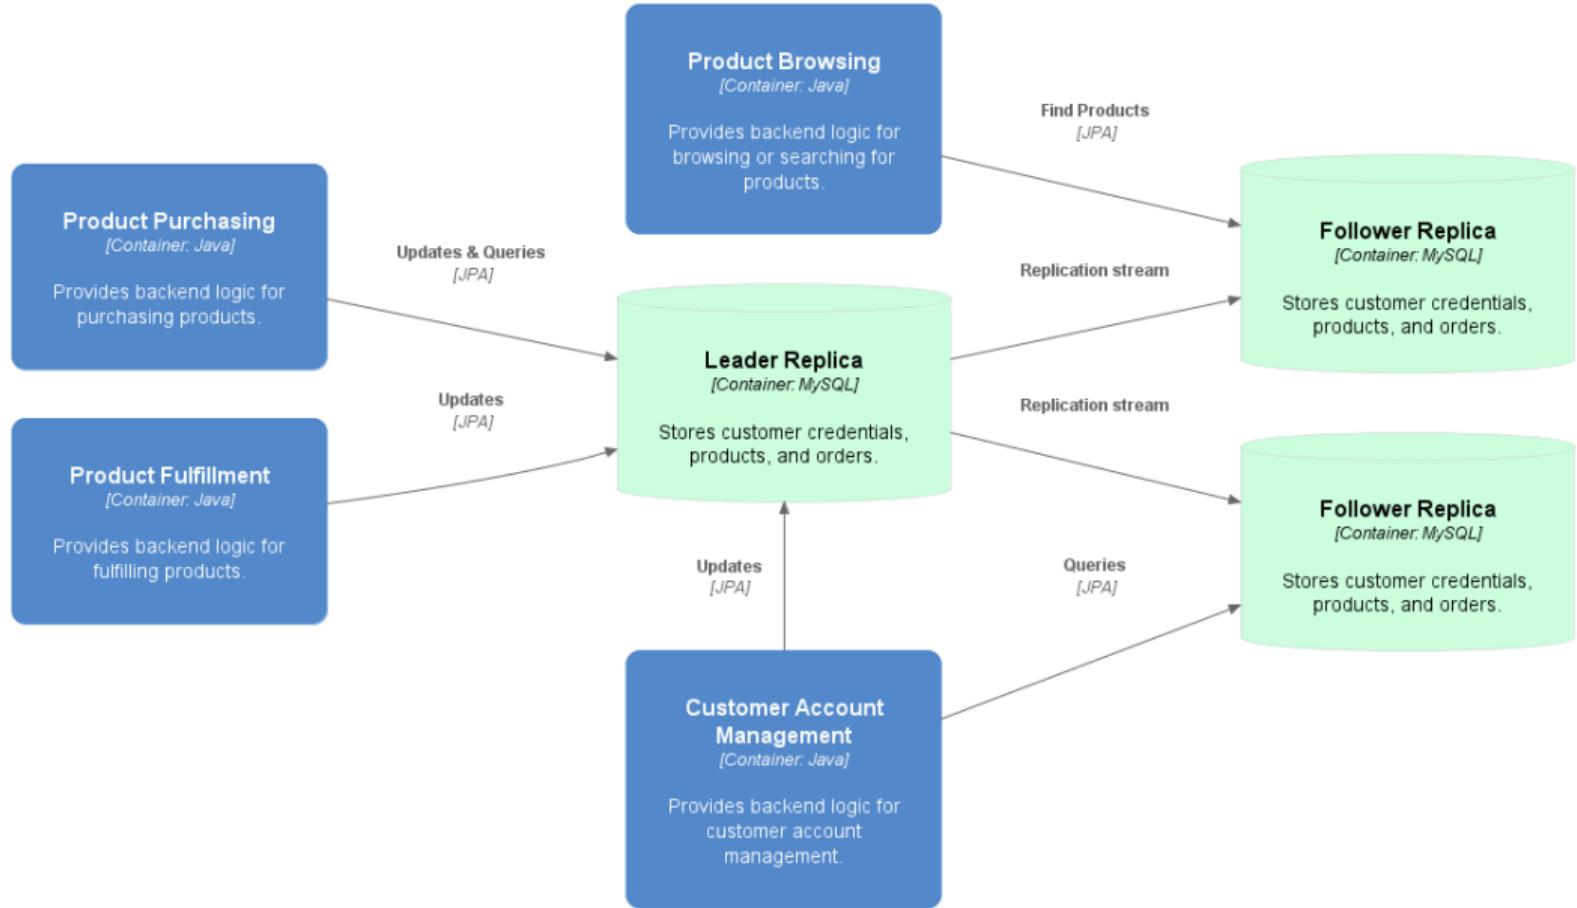
## *Leader-based Replication*

On write Writes sent to *leader*, change is propagated via change stream.

## *Leader-based Replication*

On write Writes sent to *leader*, change is propagated via change stream.

On read Any *replica* can be queried.



*Propagating Changes*

*Synchronous vs. Asynchronous*

ProductBrowsing

Leader

Follower1

1 UPDATE products SET stock=4

2 UPDATE

3 OK

4 OK

ProductBrowsing

Leader

Follower1

ProductBrowsing



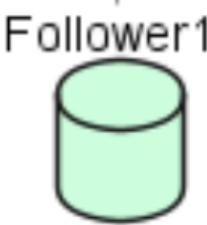
1 UPDATE products SET stock=4

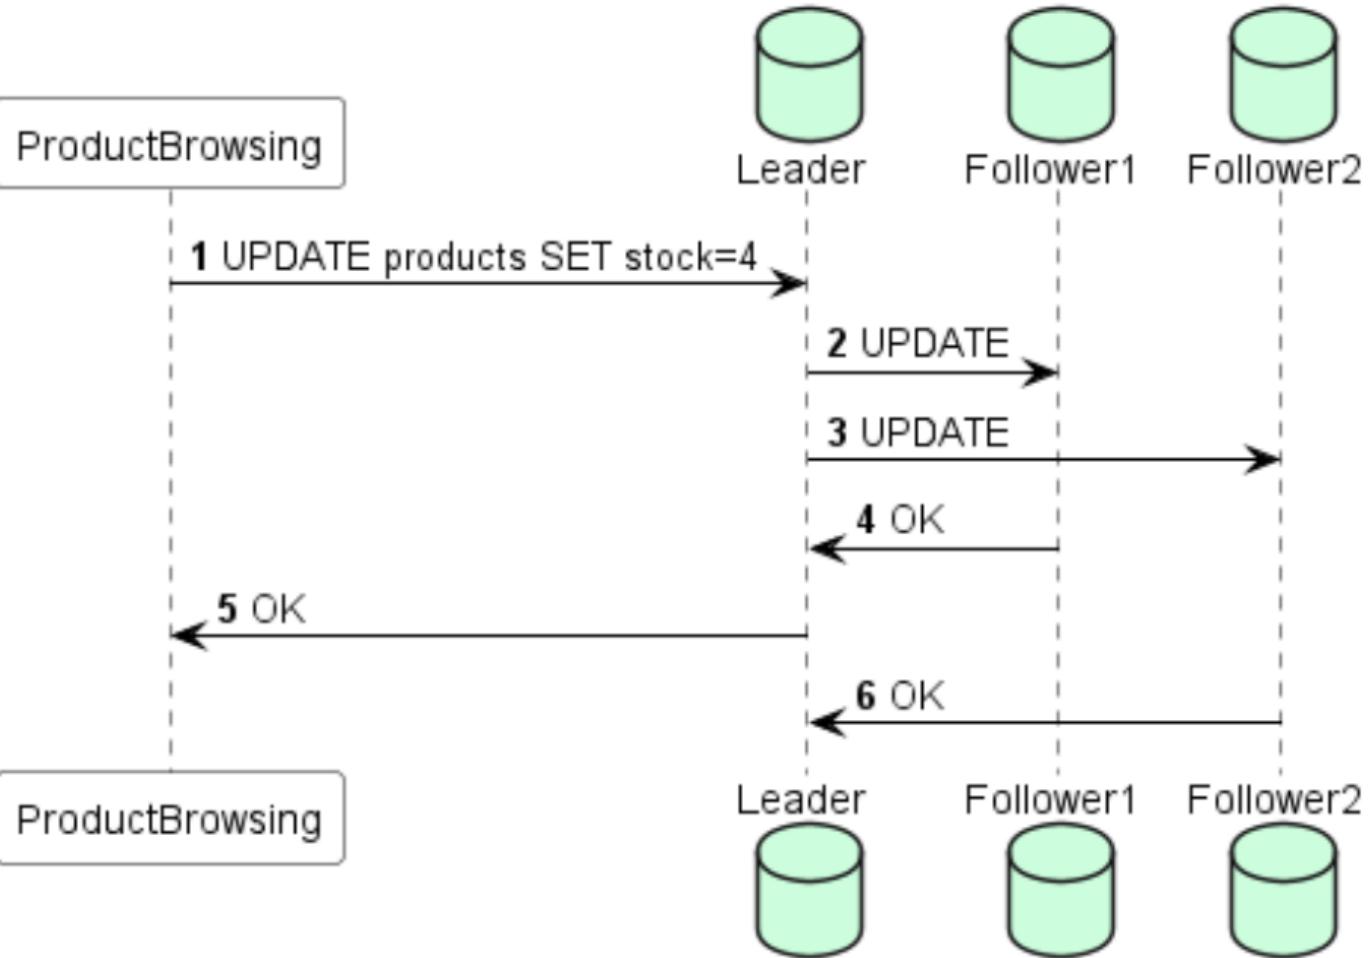
2 UPDATE

3 OK

4 OK

ProductBrowsing





## *Synchronous Propagation*

- Writes must propagate to *all followers* before being successful.

## *Synchronous Propagation*

- Writes must propagate to *all followers* before being successful.
- *Any* replica goes down, *all* replicas are un-writeable.

## *Synchronous Propagation*

- Writes must propagate to *all followers* before being successful.
- *Any* replica goes down, *all* replicas are un-writeable.
- Writes must *wait* for propagation to *all* replicas.

## *Asynchronous Propagation*

- Writes *don't* have to *wait* for propagation.

## *Asynchronous Propagation*

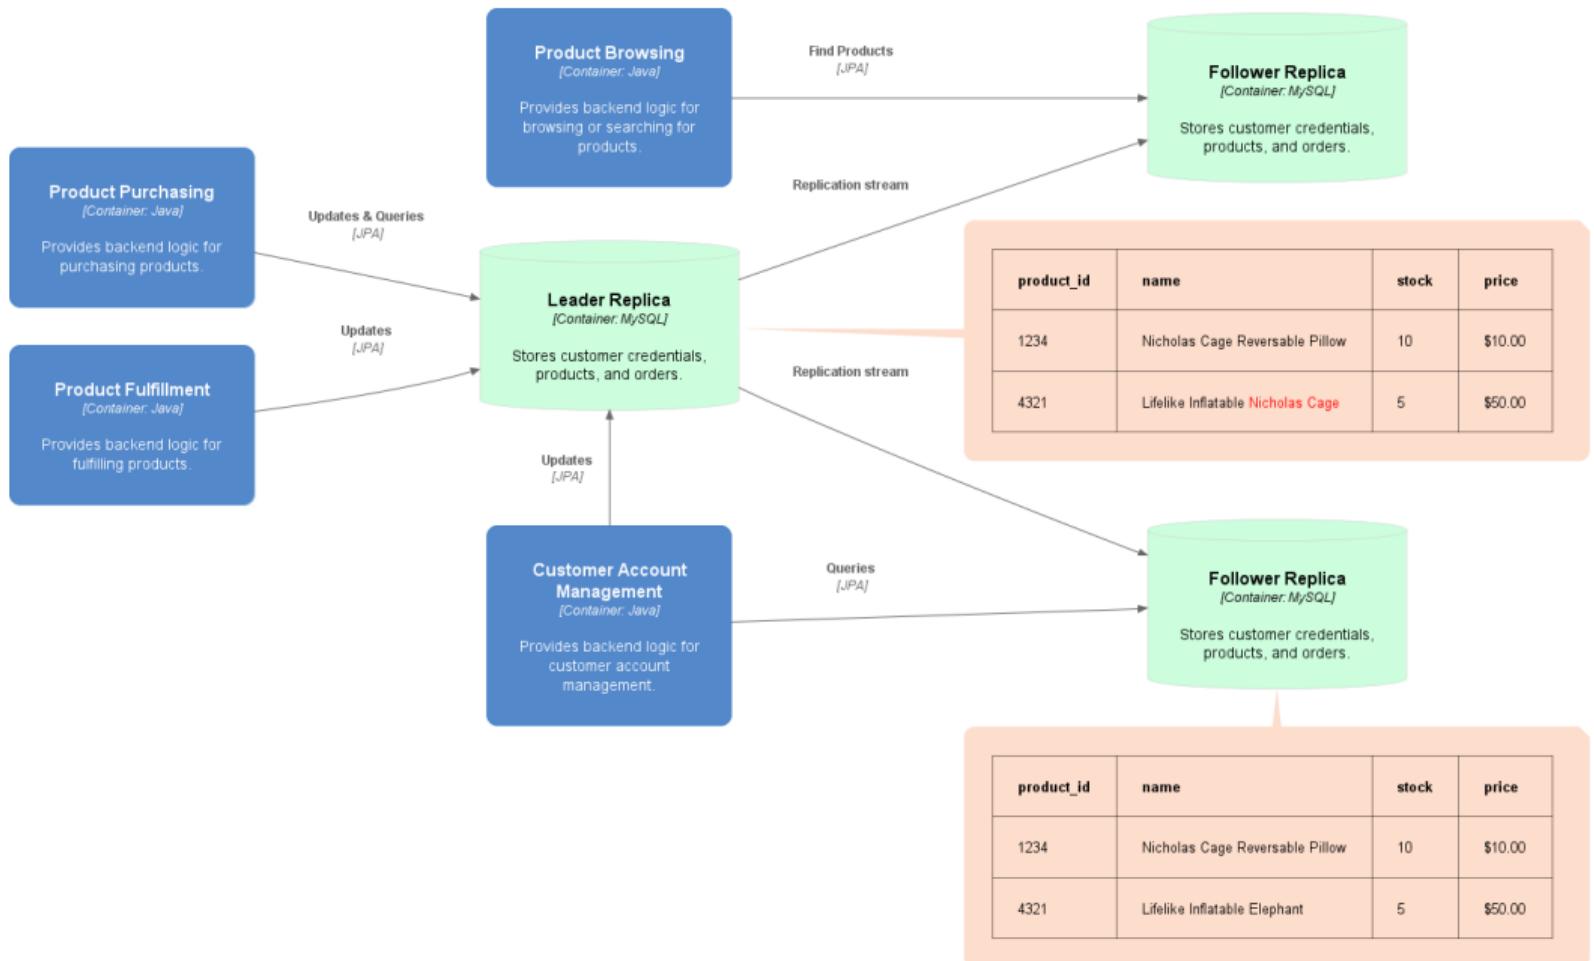
- Writes *don't* have to *wait* for propagation.
- If the leader goes down before propagating, the *write is lost*.

## *Asynchronous Propagation*

- Writes *don't* have to *wait* for propagation.
- If the leader goes down before propagating, the *write is lost*.
- Replicas can have out-dated or *stale* data.

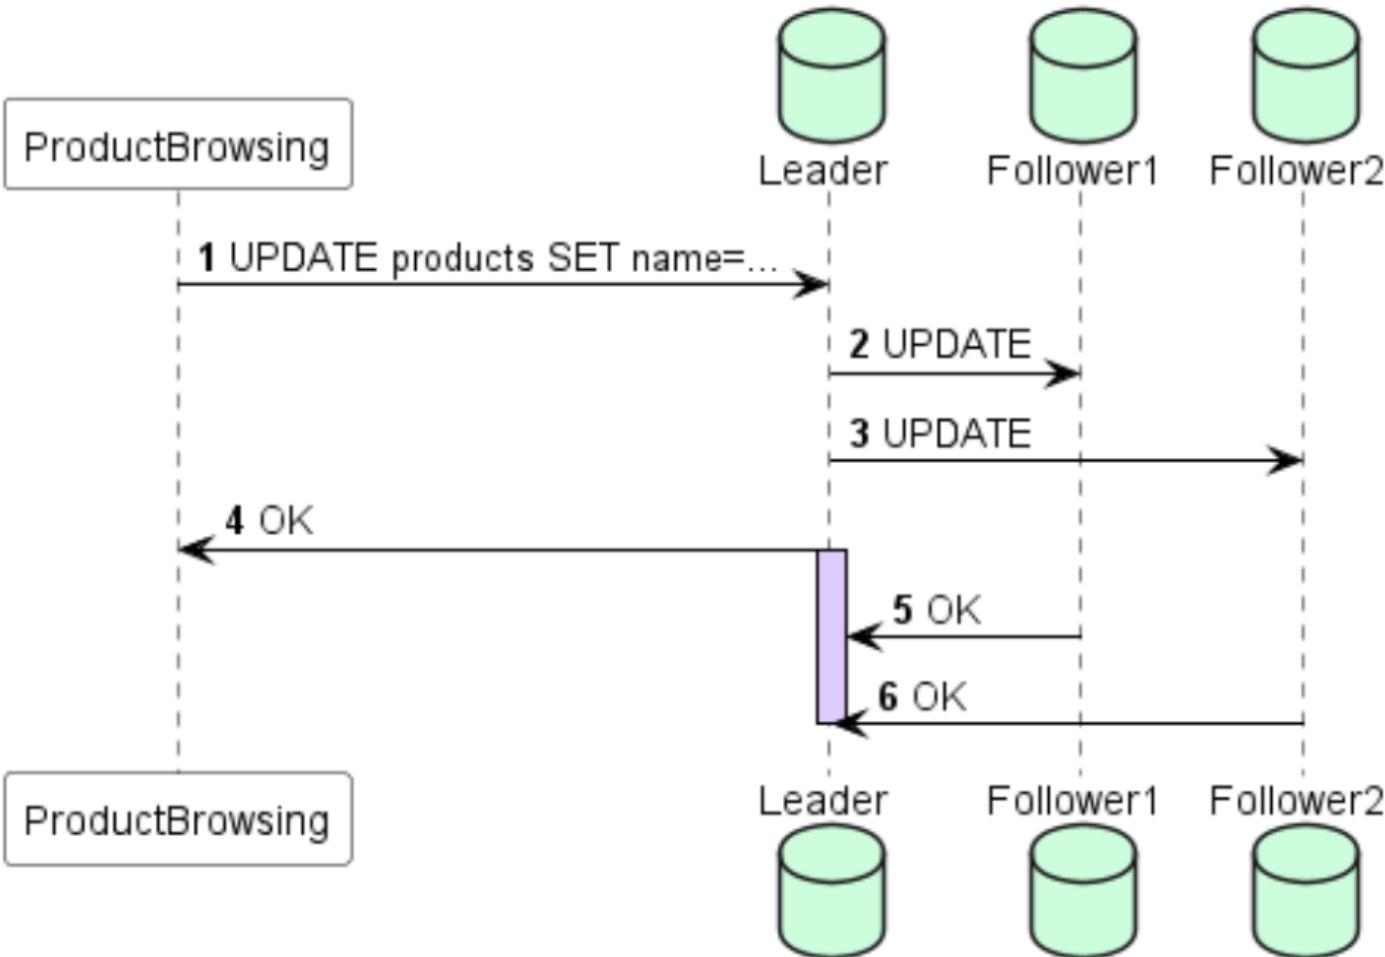
*Definition 0.* Stale data

Outdated or inconsistent data that does not reflect the latest updates, mainly due to replication delays, caching, or network issues in distributed systems.



*Definition 0.* Replication Lag

The time taken for replicas to update *stale* data.



*Eventually, all replicas must become consistent*

The system is *eventually consistent*.

*Eventual Consistency*

Sufficient? Problems?



**Brae Webb**  
@braewebb



**Brae Webb**  
@braewebb

Name: Brae

Cancel

Save

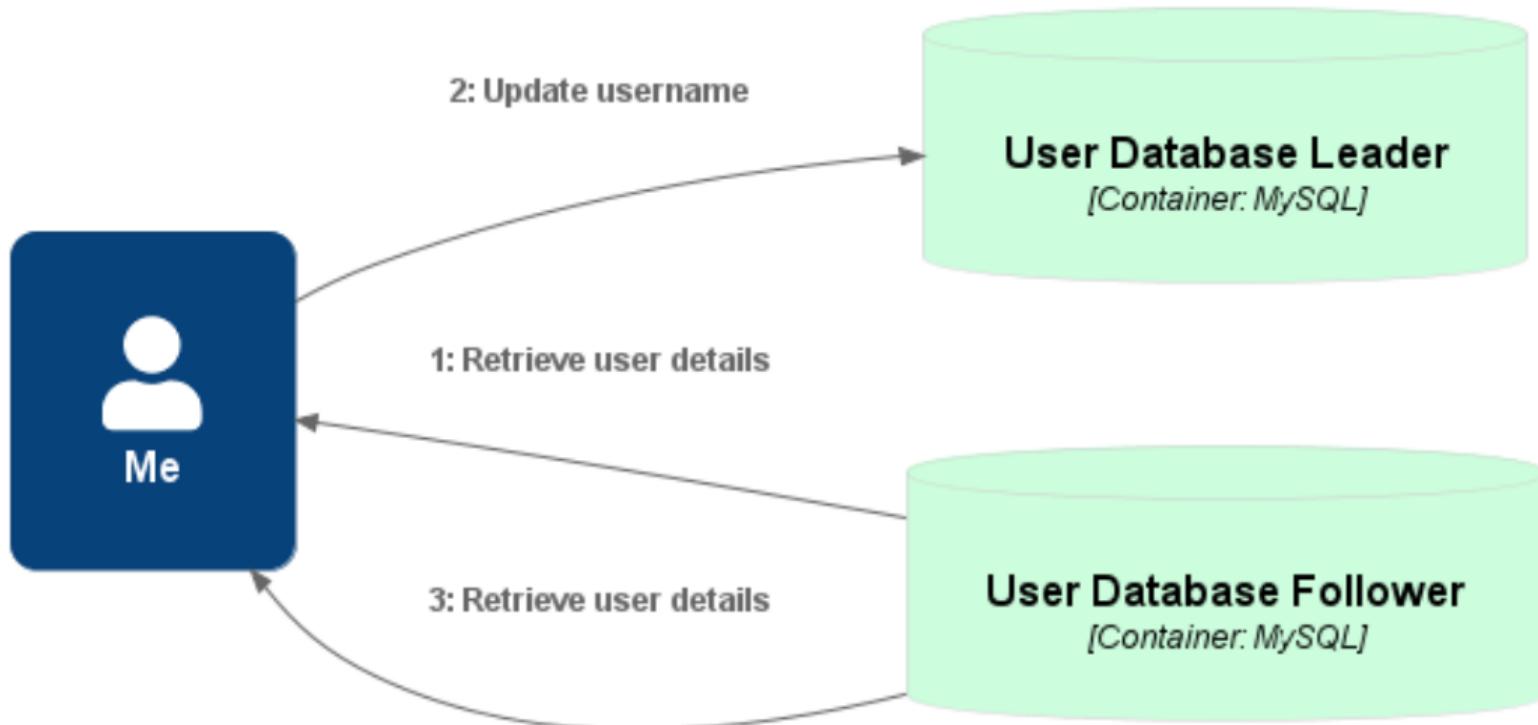


**Brae Webb**  
@braewebb

Name:	<input type="text" value="Brae"/>
<input type="button" value="Cancel"/>	<input type="button" value="Save"/>



**Brae Webb**  
@braewebb



### *Definition 0.* Read-your-writes Consistency

Users always see the updates that *they have made* (even though others see stale data).



**Brae Webb**  
@braewebb

My fist post



**Brae Webb**

@braewebb

My fist post



**Brae Webb**

@braewebb

My first post



**Brae Webb**

@braewebb

My fist post



**Brae Webb**

@braewebb

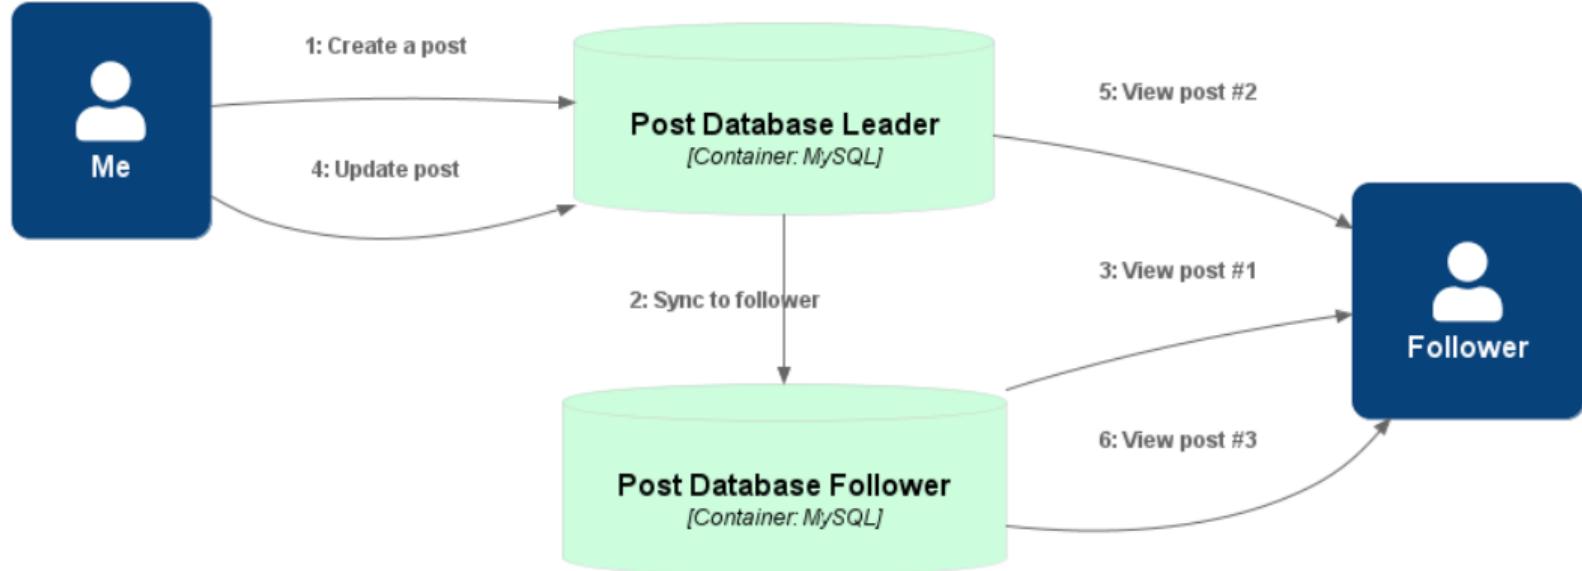
My first post



**Brae Webb**

@braewebb

My fist post



### *Definition 0.* Monotonic Reads

Once a user reads an updated value, they don't later see the old value.

### *Definition 0.* Causal Consistency

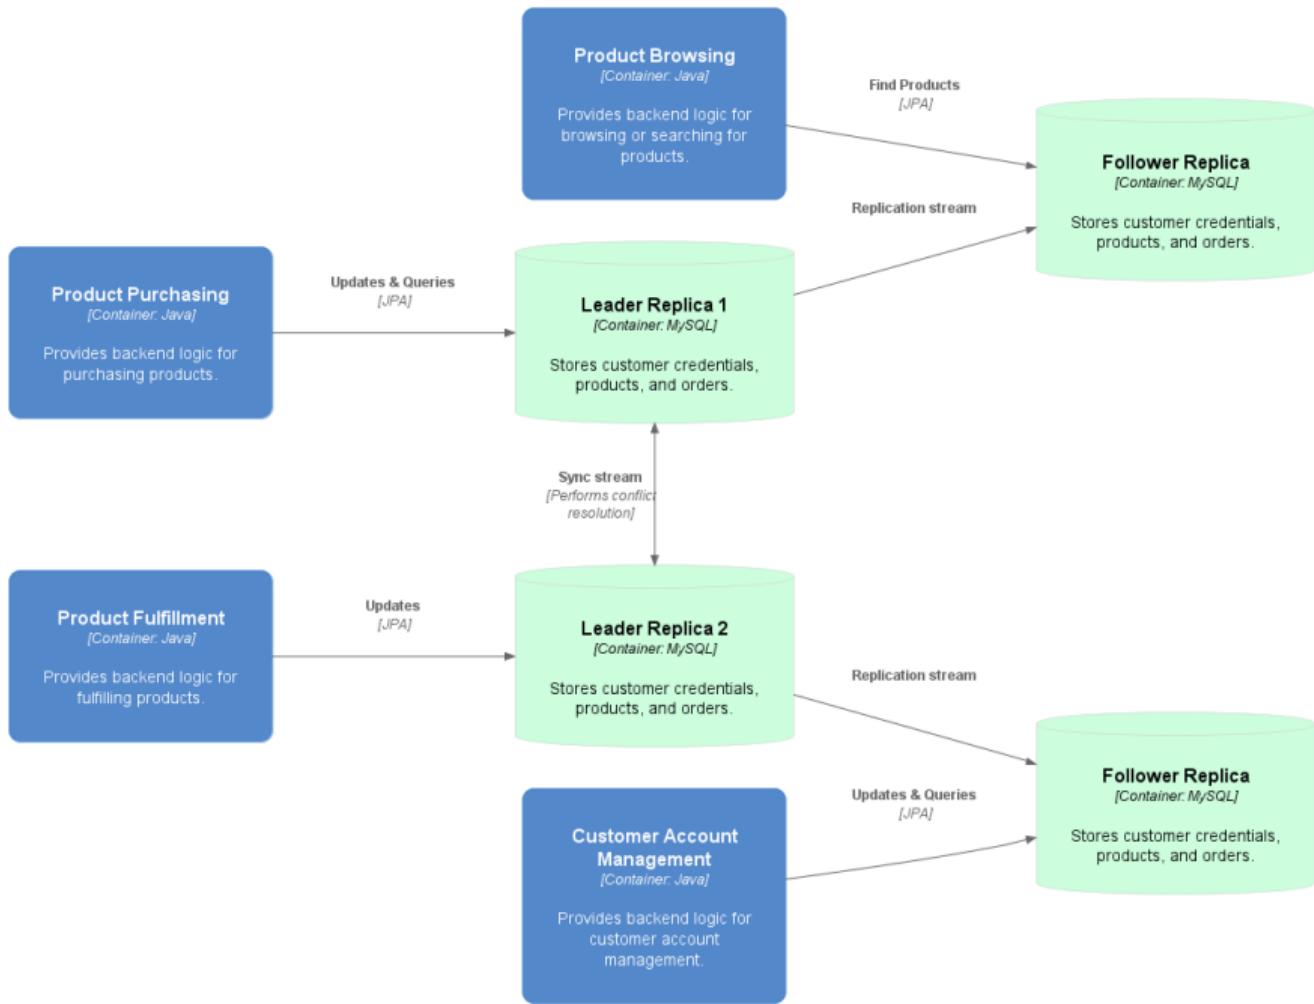
Causally related updates appear in order (e.g., comments appear under the right post).

## *Summary*

- Leader-follower databases allow *reads to scale* more effectively.
- Asynchronous propagation weakens consistency to *eventually consistent*.
- Leader-follower databases still have a *leader write bottle-neck*.

*Second approach*

Multi-leader Replication



## *Why multi-leader?*

- If you have multiple leaders, you can write to any, allowing *writes to scale*.

## *Why multi-leader?*

- If you have multiple leaders, you can write to any, allowing *writes to scale*.
- A leader going down doesn't prevent writes, giving *better fault-tolerance*.

*Question*

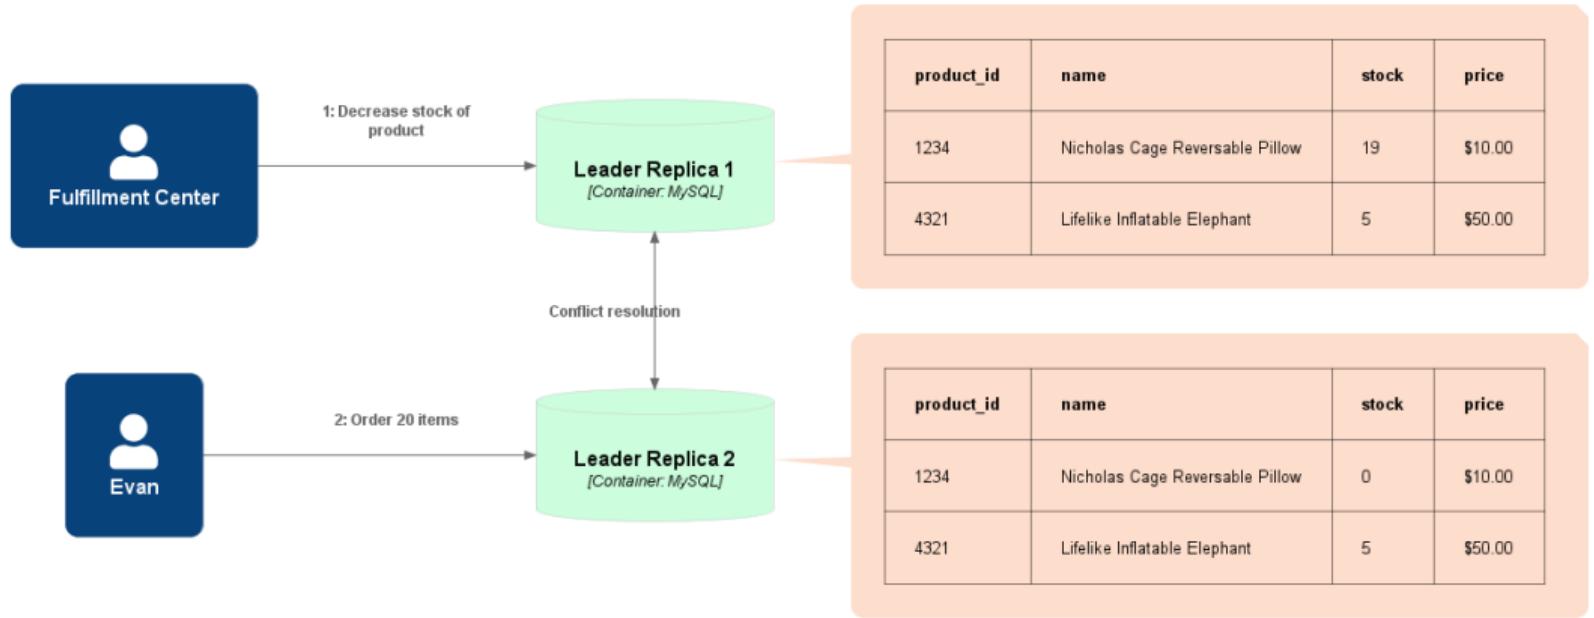
What might go wrong?

*Question*

What might go wrong?

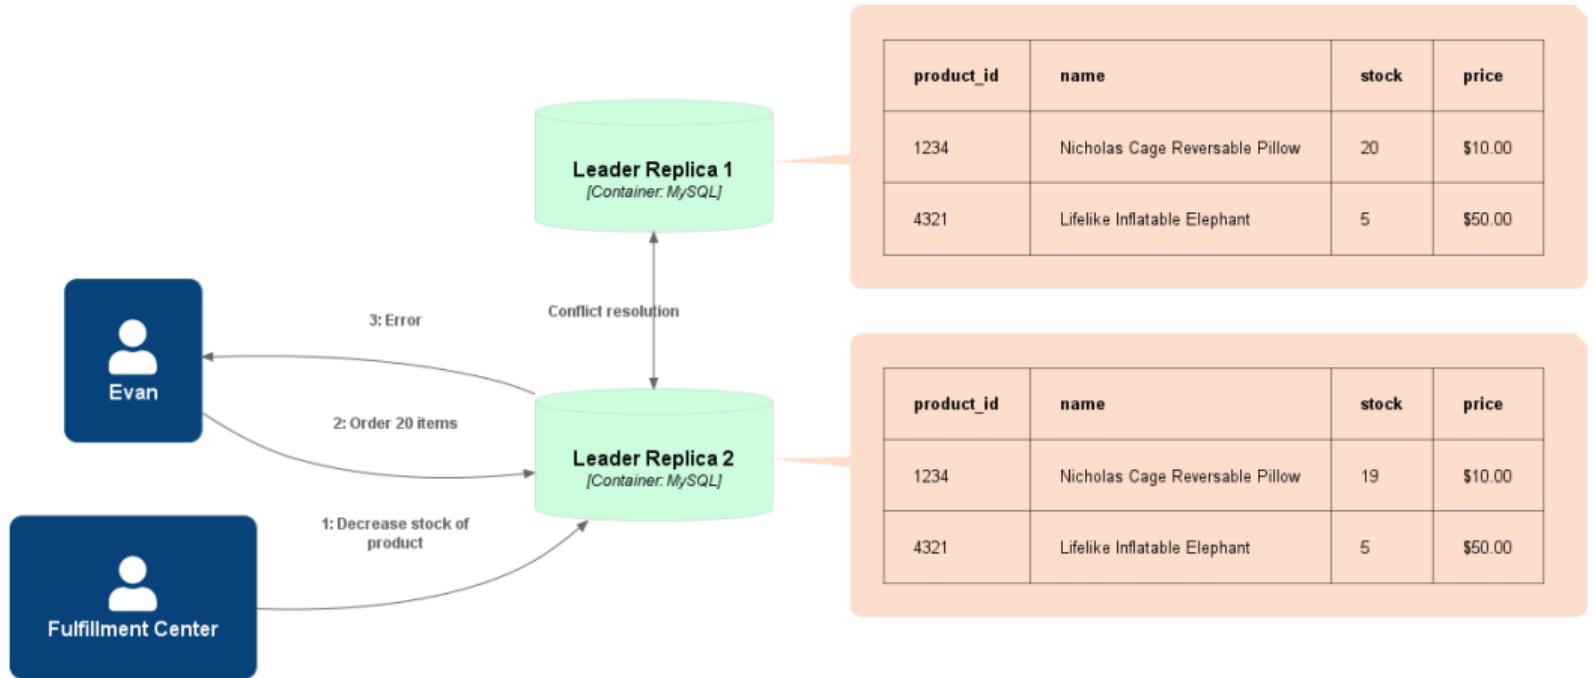
*Answer*

Write conflicts



*Where possible*

Avoid write conflicts



*Where impossible*

Convergence

## *Convergence Strategies*

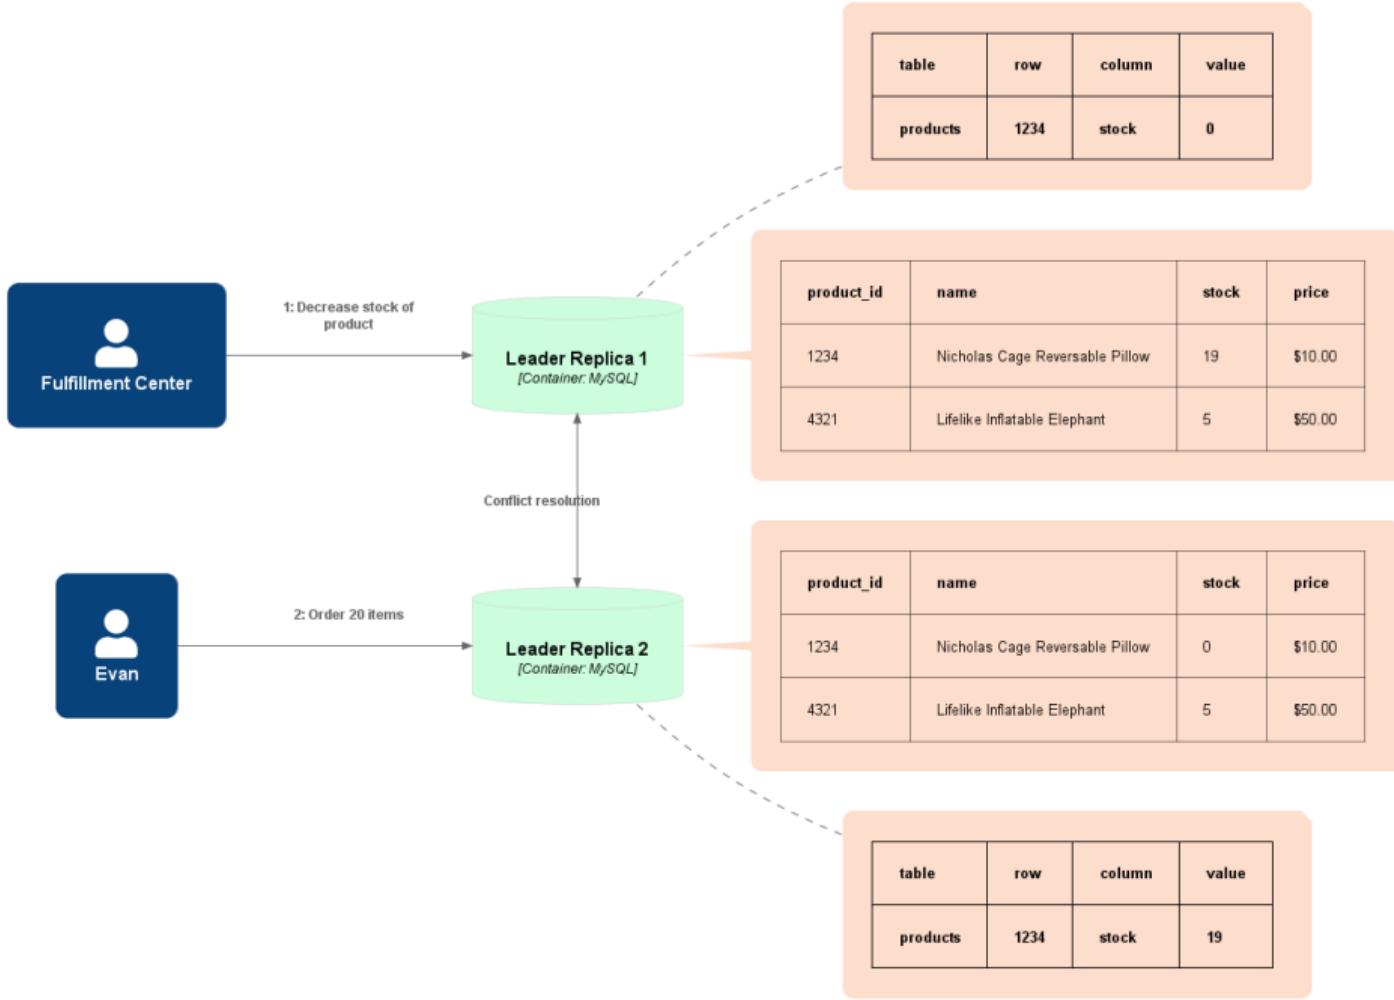
- Assign each *write* a unique ID

## *Convergence Strategies*

- Assign each *write* a unique ID
- Assign each *leader replica* a unique ID

## *Convergence Strategies*

- Assign each *write* a unique ID
- Assign each *leader replica* a unique ID
- Custom resolution logic



## *Resolving Conflicts*

**On Write** When a conflict is first noticed, take proactive resolution action.

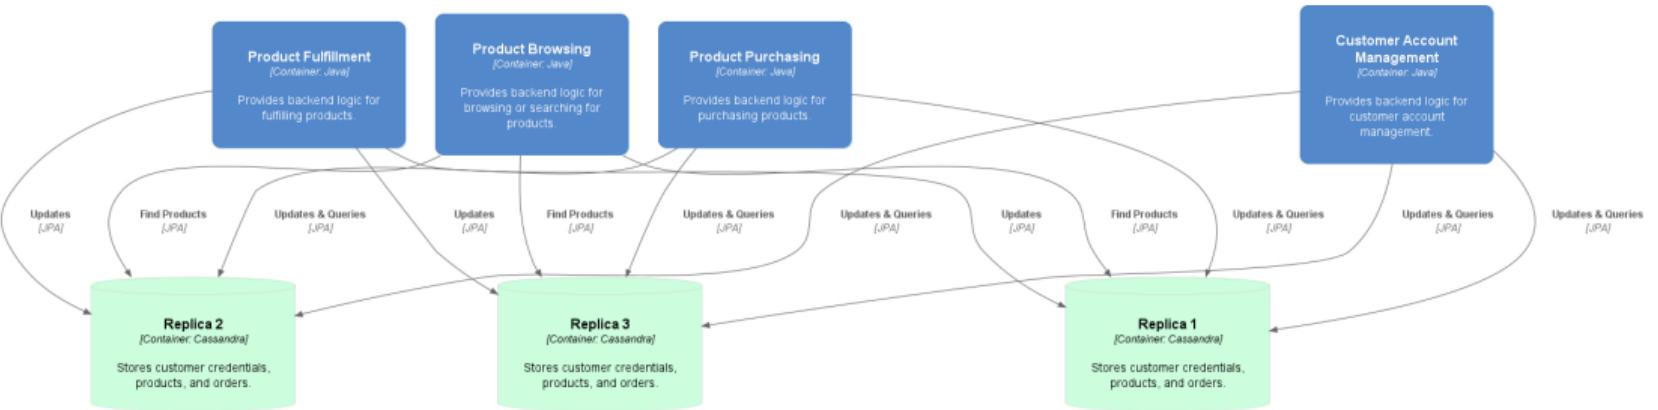
Example: Last Write Wins (LWW) in DynamoDB.

**On Read** Stores all conflicting versions on write. When a conflict is next read, ask for a resolution.

Example: *git pull*

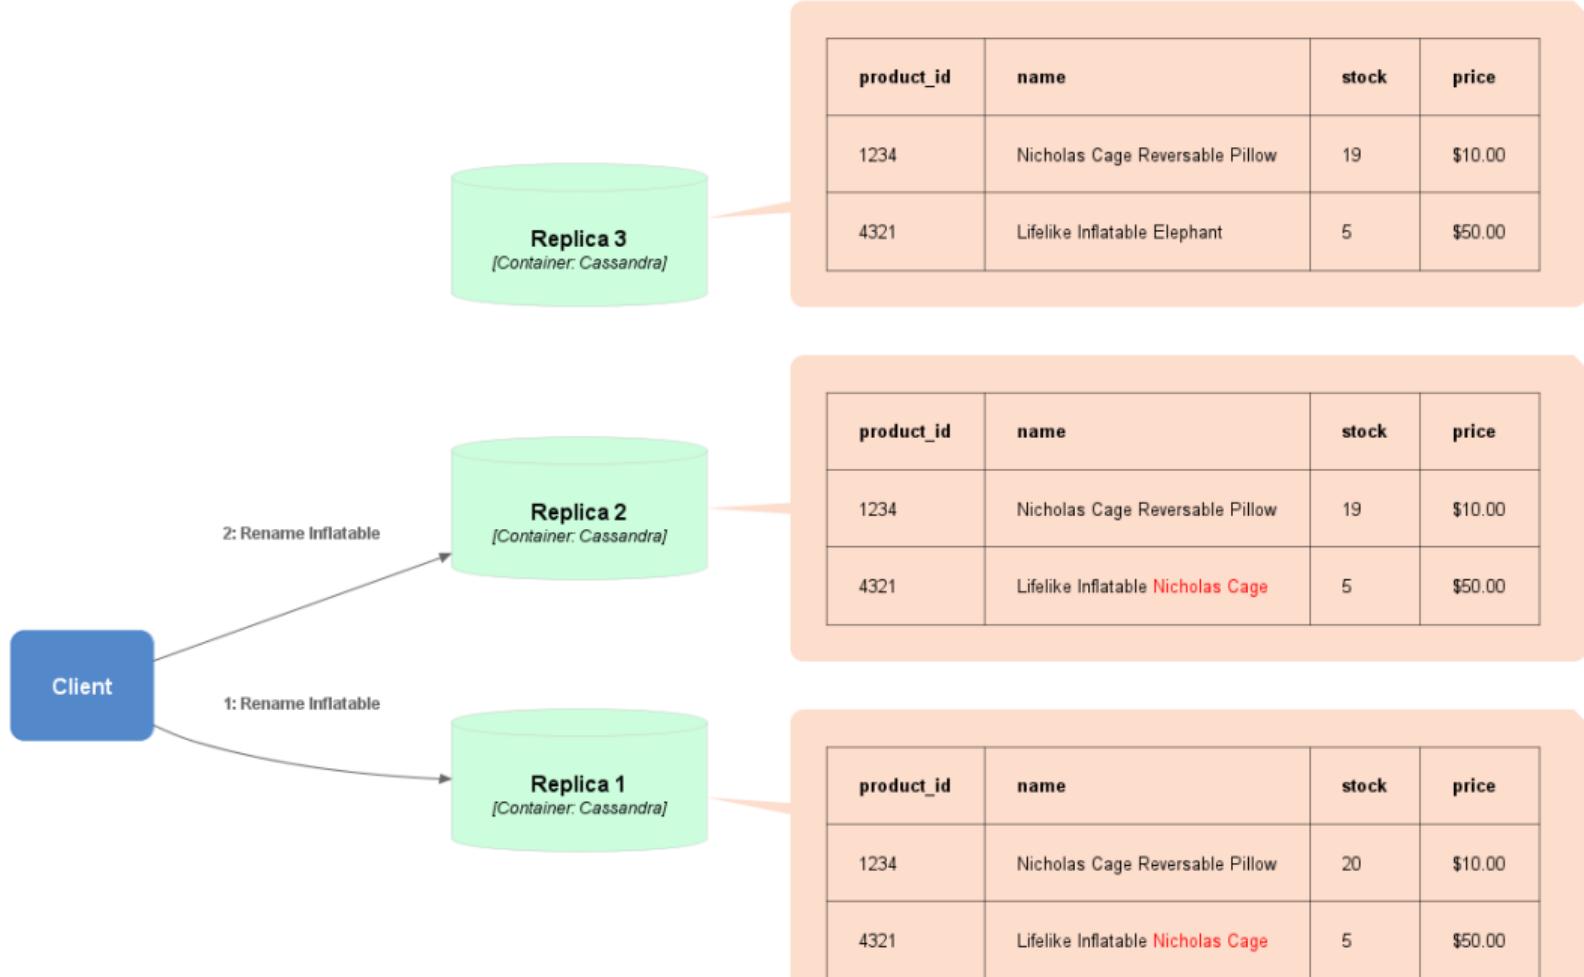
*Third Approach*

Leaderless Replication



*How do they work?*

Each read/write is sent to *multiple* replicas.





*How are changes propagated?*

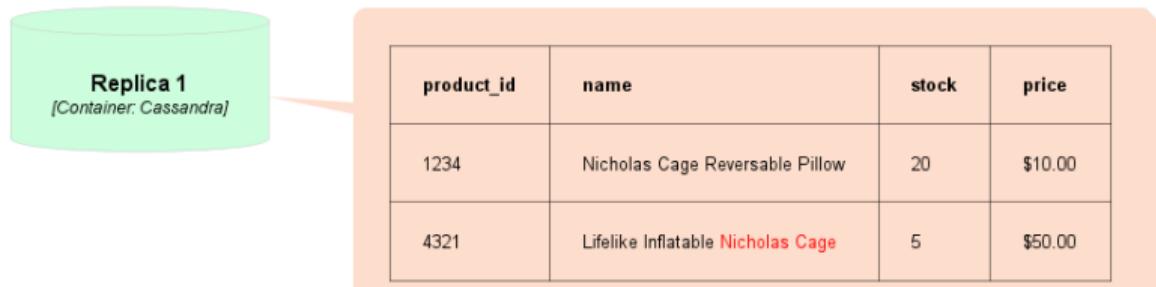
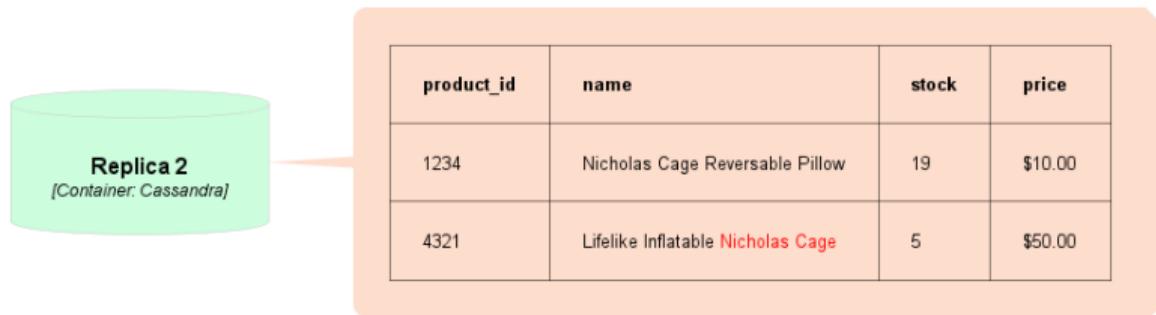
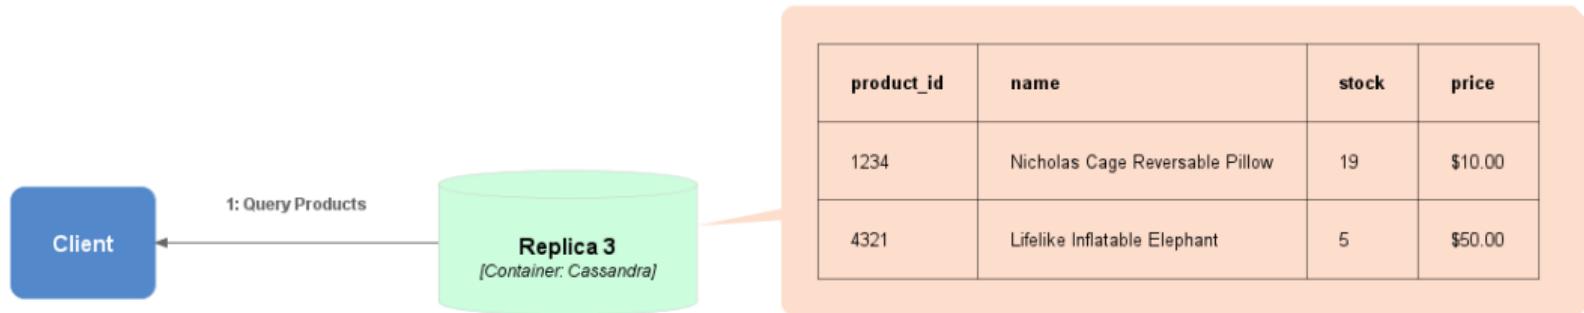
- Read Repair

*How are changes propagated?*

- Read Repair
- Anti-Entropy Process

*Question*

How do we know it's consistent?



*Question*

How do we know it's consistent?

*Question*

How do we know it's consistent?

*Answer*

Quorum Reads and Writes

## *Quorum Consistency*

$$w + r > n$$

*n* total replicas

*w* amount of replicas to *write* to

*r* amount of replicas to *read* from

## *Quorum Consistency*

$$2 + 2 > 3$$

*n* total replicas

*w* amount of replicas to *write* to

*r* amount of replicas to *read* from

## *Quorum Consistency*

$$1 + 3 > 3$$

*n* total replicas

*w* amount of replicas to *write* to

*r* amount of replicas to *read* from



$n$  total replicas

$w$  amount of replicas to *write* to

$r$  amount of replicas to *read* from

*Question*

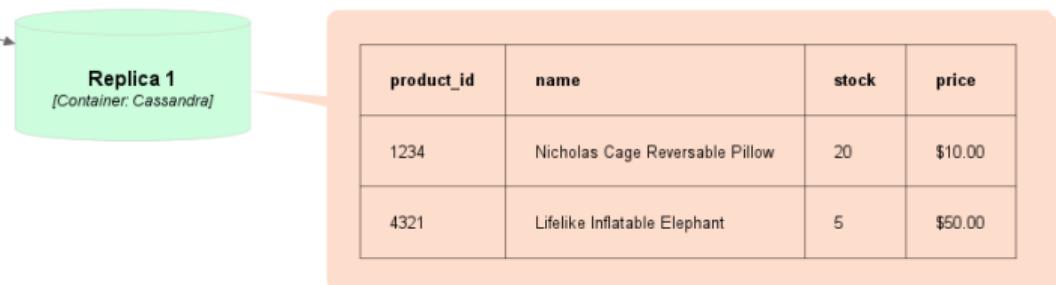
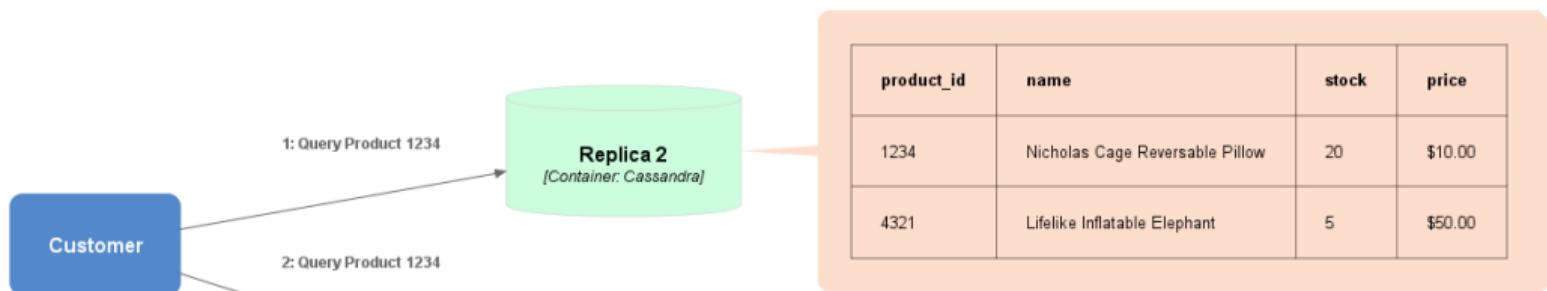
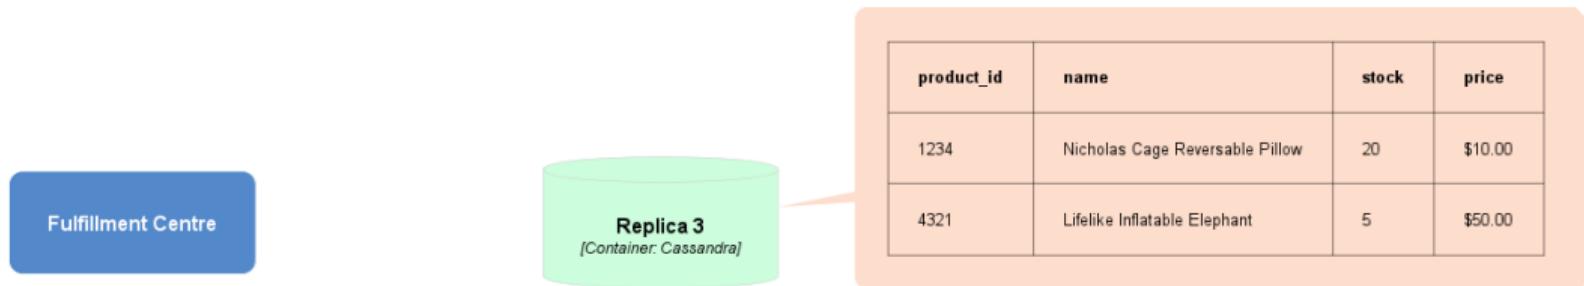
What about write conflicts?

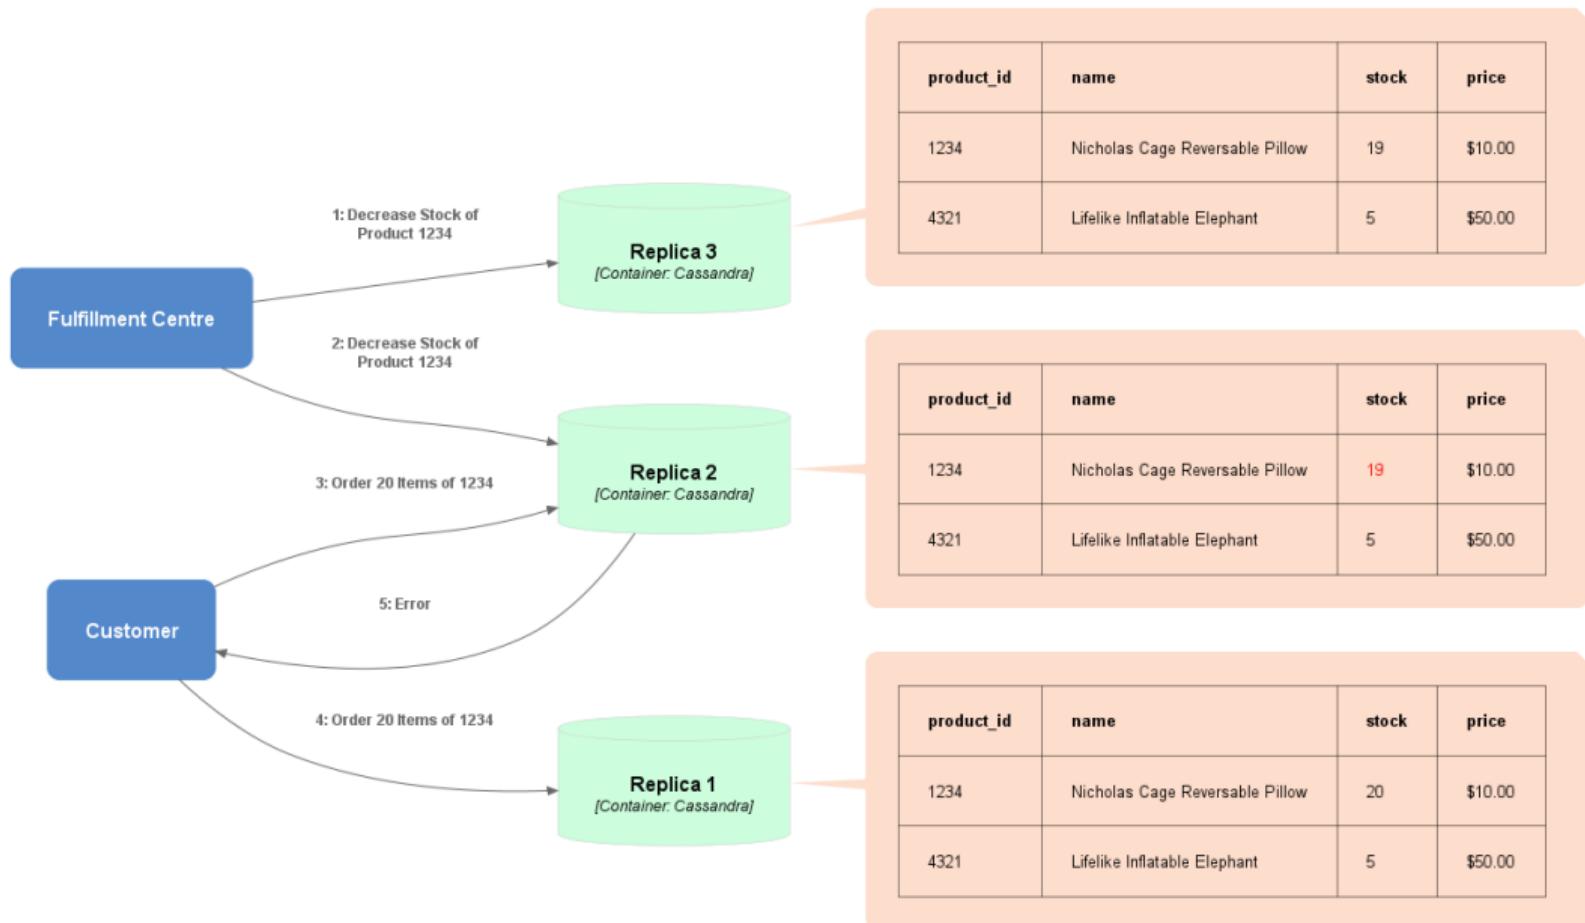
*Question*

What about write conflicts?

*Answer*

Same problem as with Multi-leader replication.





## *Summary*

- *Replication* copies data to multiple replicas.

## *Summary*

- *Replication* copies data to multiple replicas.
- *Leader-based* replication is most common and simplest.

## *Summary*

- *Replication* copies data to multiple replicas.
- *Leader-based* replication is most common and simplest.
- Replication introduces *eventual consistency*.

## *Summary*

- *Replication* copies data to multiple replicas.
- *Leader-based* replication is most common and simplest.
- Replication introduces *eventual consistency*.
- *Multi-leader* replication scales writes as well as reads but introduces *write conflicts*.

## *Summary*

- *Replication* copies data to multiple replicas.
- *Leader-based* replication is most common and simplest.
- Replication introduces *eventual consistency*.
- *Multi-leader* replication scales writes as well as reads but introduces *write conflicts*.
- *Leaderless* replication is another approach which keeps the problems of multi-leader.

*Question*

How do we fix database scaling issues?

*Question*

How do we fix database scaling issues?

*Answer*

- *Replication*
- Partitioning
- Independent databases

*Question*

How do we fix database scaling issues?

*Answer*

- Replication
- *Partitioning*
- Independent databases

## *Definition 0.* Partitioning

Split the data of a system onto multiple nodes.

These nodes are *partitions*.



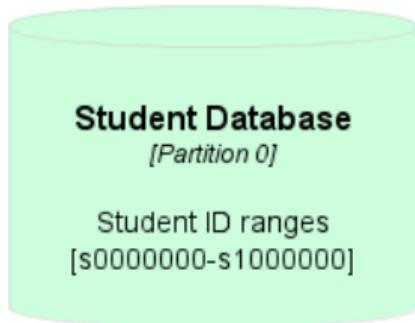
product_id	name	stock	price
4321	Lifelike Elephant Inflatable	5	\$50.00



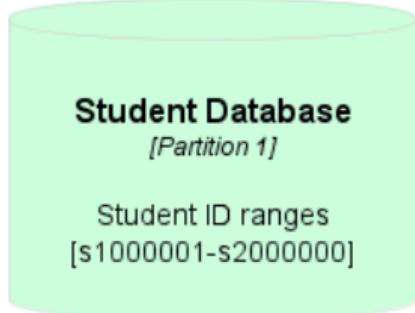
product_id	name	stock	price
1234	Nicholas Cage Reversible Pillow	10	\$10.00

*Question*

How should we decide which data is stored where?



student_id	name	...
s0746283	Bobby Tables	...
...	...	...



student_id	name	...
s1637285	Brae Webb	...
...	...	...

*Question*

What is the problem with this?

*Question*

What is the problem with this?

*Answer*

Over time some partitions become inactive,  
while others receive almost all load.

*Question*

How should we decide where data is stored?

*Question*

How should we decide where data is stored?

*Answer*

Maximize spread of requests, avoiding *skewing*.

*Question*

Have we seen this before?

*Question*

Have we seen this before?

*Answer*

Hashing?

*Question*

What is the problem with this?

*Question*

What is the problem with this?

*Answer*

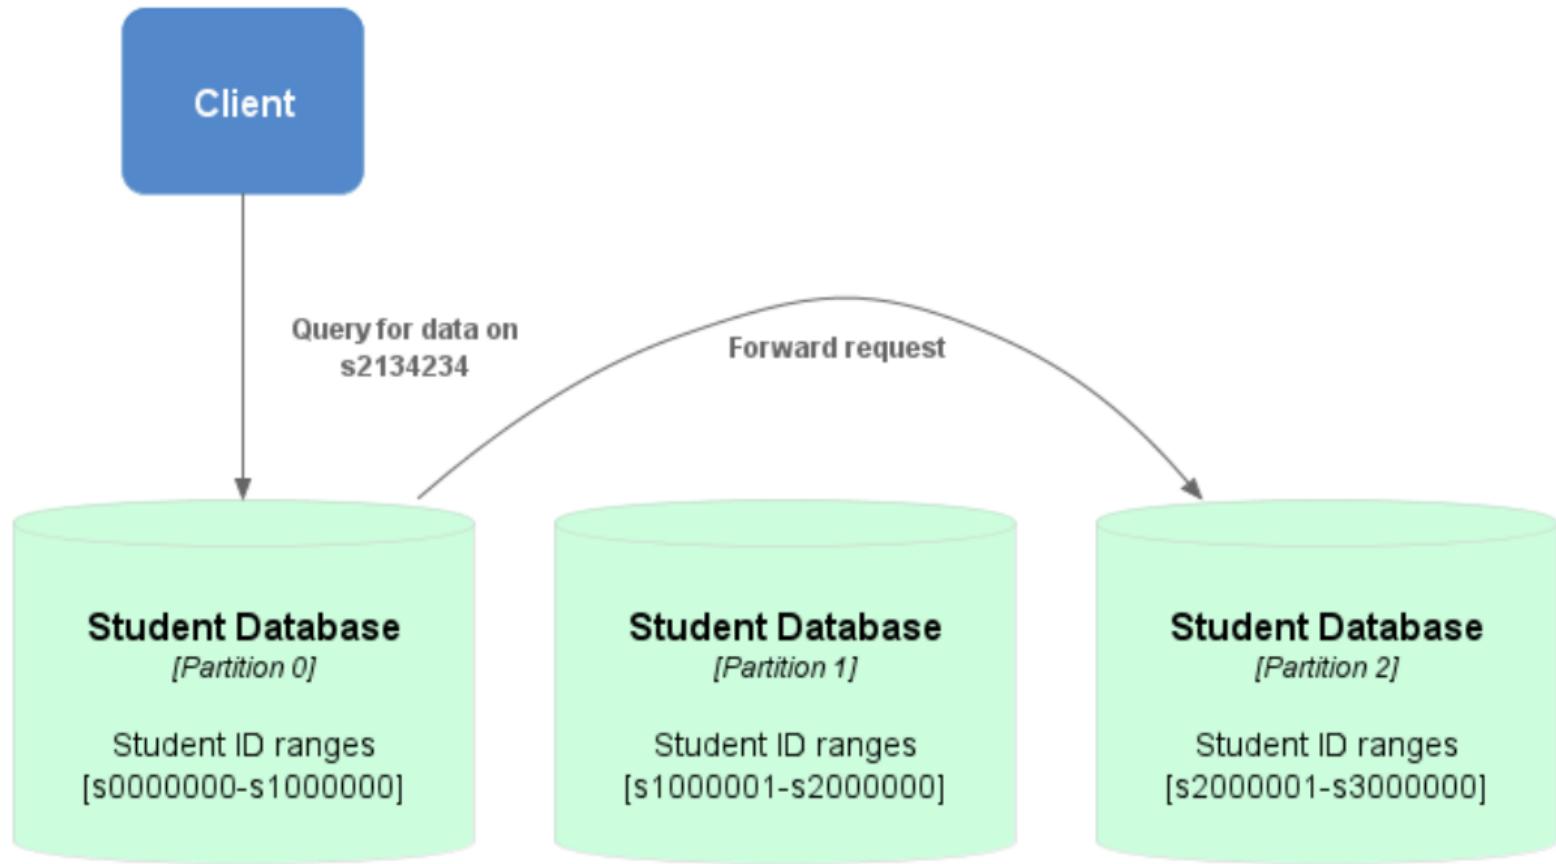
Range queries are inefficient, i.e. get all students between s4444444 and s4565656.

*Question*

How do we route queries?

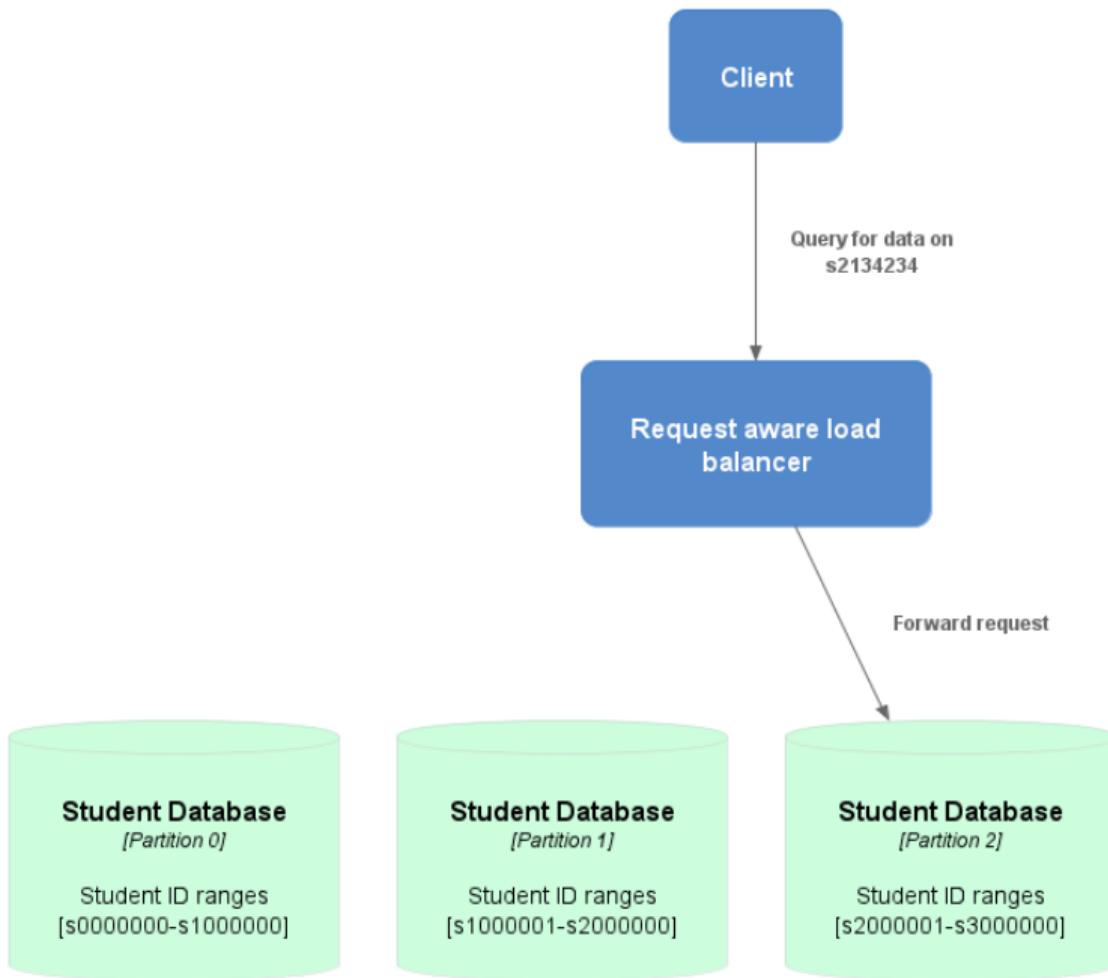
## *Query-Insensitive Load Balancer*

Randomly route to any node, responsibility of the node to re-route to the correct node.



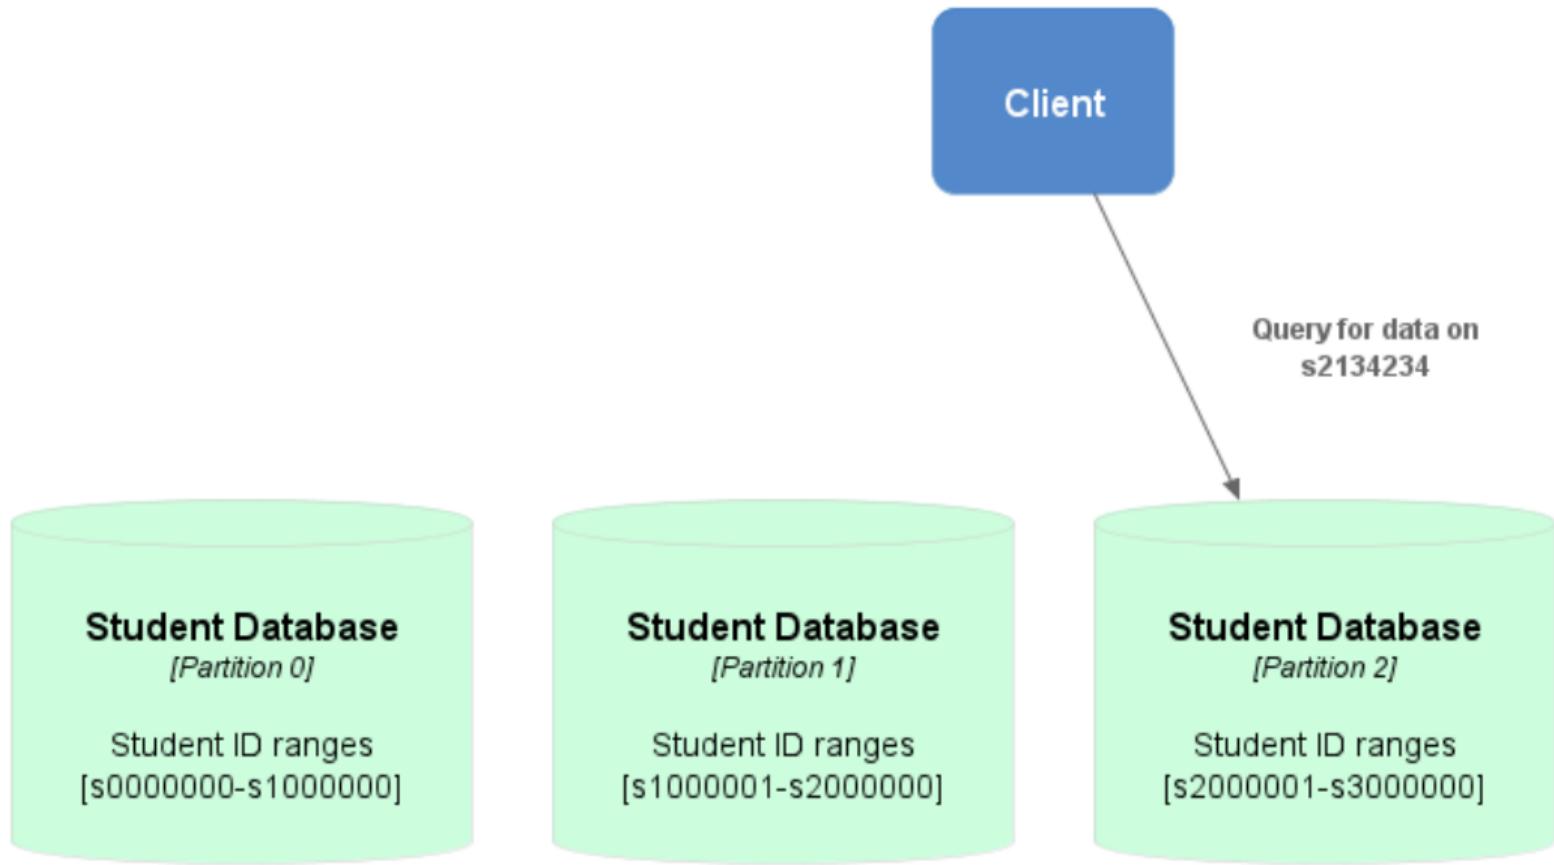
## *Query-Sensitive Load Balancer*

Load balancer understands which queries should be forwarded to which node.



## *Client-aware Queries*

Place the responsibility on clients to choose the correct node.



## *Summary*

- *Partitioning* splits data across multiple nodes.

## *Summary*

- *Partitioning* splits data across multiple nodes.
- Requires a *consistent method* to choose appropriate node.

## *Summary*

- *Partitioning* splits data across multiple nodes.
- Requires a *consistent method* to choose appropriate node.
- Partitioning by *primary key* can create *skewing*.

## *Summary*

- *Partitioning* splits data across multiple nodes.
- Requires a *consistent method* to choose appropriate node.
- Partitioning by *primary key* can create *skewing*.
- Partitioning by *hash* makes range queries less efficient.

## *Summary*

- *Partitioning* splits data across multiple nodes.
- Requires a *consistent method* to choose appropriate node.
- Partitioning by *primary key* can create *skewing*.
- Partitioning by *hash* makes range queries less efficient.
- Three approaches to *routing requests*.

### *Disclaimer*

We have ignored the *hard* parts of replication.

*Question*

How do we fix database scaling issues?

*Question*

How do we fix database scaling issues?

*Answer*

- Replication
- *Partitioning*
- Independent databases

*Question*

How do we fix database scaling issues?

*Answer*

- Replication
- Partitioning
- *Independent databases*

## *Summary*

- Replications

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless
  - Eventual consistency

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless
  - Eventual consistency
  - Write conflicts

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless
  - Eventual consistency
  - Write conflicts
- Partitioning

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless
  - Eventual consistency
  - Write conflicts
- Partitioning
  - Consistent method to pick nodes for data

## *Summary*

- Replications
  - Leader-based, multi-leader, and leaderless
  - Eventual consistency
  - Write conflicts
- Partitioning
  - Consistent method to pick nodes for data
  - Avoiding skewing