Software at Scale

Software Architecture

Brae Webb

March 27, 2023

How many concurrent users can your software handle?

How many concurrent users can your software handle?

Maybe *400*? Maximum.

Question

How many concurrent users can your software handle?

Answer

Maybe 400[1]? Maximum.

---

[1] HTTP server on a t2.micro EC2 instance

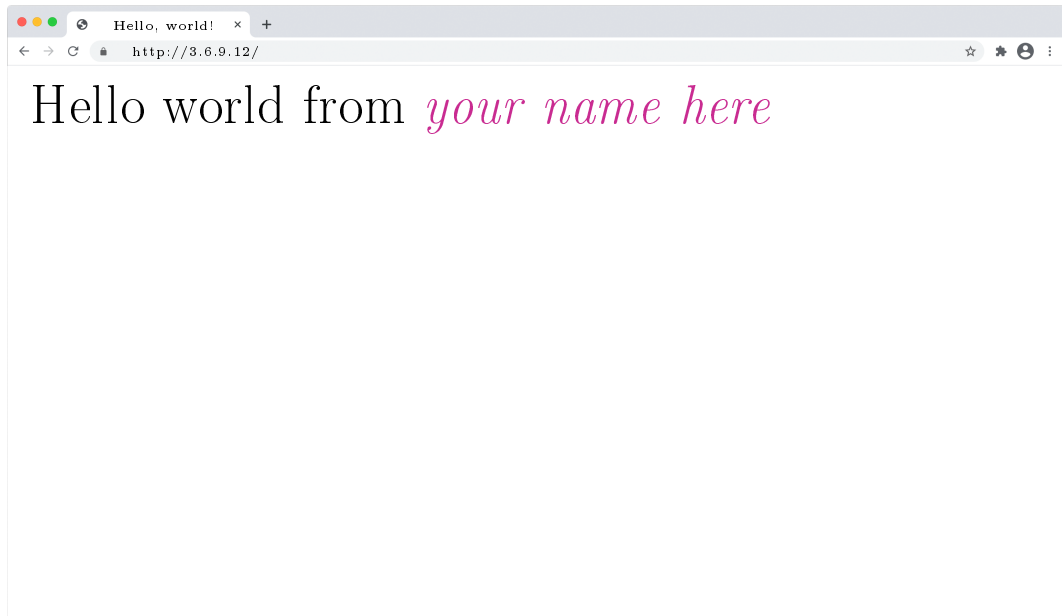**Definition 1.** Stress Testing

Measure the robustness of software by pushing usage to an extreme.

Let's build 'hello world'

Hello world from *your name here*

Hello, world!

http://3.6.9.12/

Hello world from Brae

```
» cat hello-server.tf

resource "aws_instance" "hello-server" {
    ami = "ami-04902260ca3d33422"
    instance_type = "t2.micro"
}
```

```
» cat hello-server.tf
```

```
1  resource "aws_instance" "hello-server" {
2      ami = "ami-04902260ca3d33422"
3      instance_type = "t2.micro"
4
5      user_data = file("./setup.sh")
6  }
```

```
» cat setup.sh

#!/bin/bash
yum -y install httpd
systemctl enable httpd
systemctl start httpd
echo '<html><title>Hello, world!</title><h1>Hello world from Brae</h1></html>' > /
    var/www/html/index.html
```

```
» cat hello-server.tf
```

```
1  resource "aws_instance" "hello-server" {
2      ami = "ami-04902260ca3d33422"
3      instance_type = "t2.micro"

5      user_data = file("./setup.sh")

7      security_groups = [
8          aws_security_group.hello-server.name
9      ]
10 }
```

```
» cat hello-server.tf

resource "aws_instance" "hello-server" {
    ami = "ami-04902260ca3d33422"
    instance_type = "t2.micro"

    user_data = file("./setup.sh")

    security_groups = [
        aws_security_group.hello-server.name
    ]

    tags = {
        Name = "hello-server"
    }
}
```

Starting the server

```
1  >> terraform init
2  >> terraform plan
3  >> terraform apply
```

Before

3.6.9.12

http://3.6.9.12/

This site can't be reached

3.6.9.12 took too long to respond.

Hello, world!

http://3.6.9.12/

Hello world from Brae
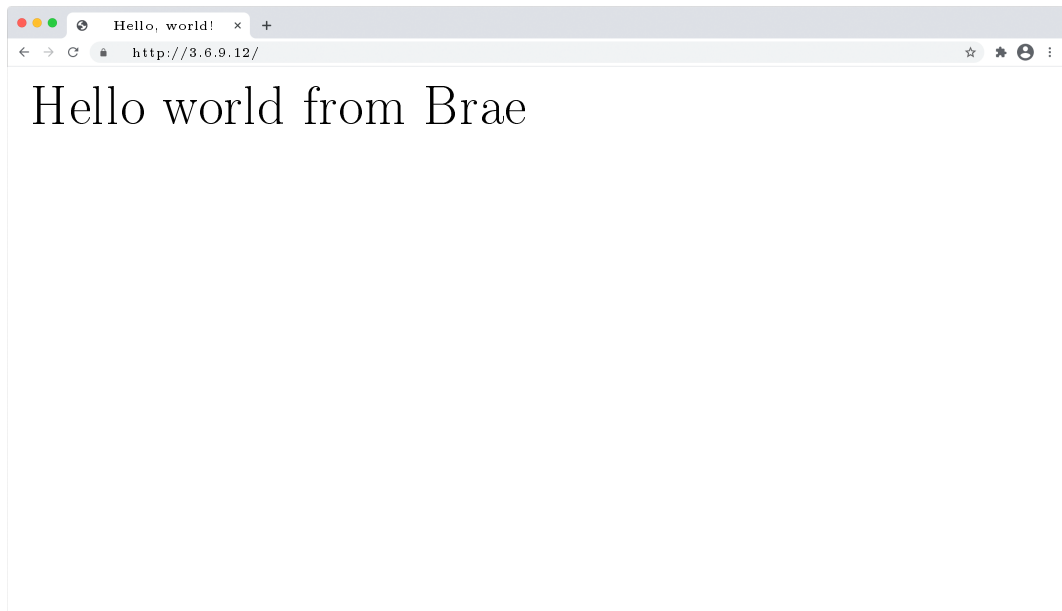
How much traffic can this website handle?

```
» cat stress-test.js

1  import http from 'k6/http';
2  import { check, sleep } from 'k6';
3
4  const IP = "http://3.6.9.12/";
5  export default function() {
6      const res = http.get(IP);
7      check(res, { 'status was 200': (r) => r.status == 200 });
8      sleep(1);
9  }
```

```
» cat stress-test.js

import http from 'k6/http';
import { check, sleep } from 'k6';

const IP = "http://3.6.9.12/";
export const options = {
    stages: [
        { duration: '2m', target: 100 },
    ],
};
export default function() {
    const res = http.get(IP);
    check(res, { 'status was 200': (r) => r.status == 200 });
    sleep(1);
}
```

Run the tests

```
>> k6 run stress-test.js
```

Looks good so far

```
1    status was 200
2      100% - 347867 / 0

4  checks.........................: 100%
5  data_received..................: 100 MB 44 kB/s
6  data_sent......................: 27 MB 12 kB/s
7  iterations.....................: 347997 152.552084/s
8  vus............................: 1 min=1 max=400
```

Let's upgrade the traffic

```
» cat stress-test.js
```

```js
1  export const options = {
2      stages: [
3          { duration: '2m', target: 100 },
4          { duration: '5m', target: 100 },
5          { duration: '2m', target: 200 },
6          { duration: '5m', target: 200 },
7          { duration: '2m', target: 300 }, // around the breaking point
8          { duration: '5m', target: 300 },
9          { duration: '2m', target: 400 }, // beyond the breaking point
10         { duration: '5m', target: 400 },
11         { duration: '2m', target: 0 }, // scale down
12     ],
13 };
```

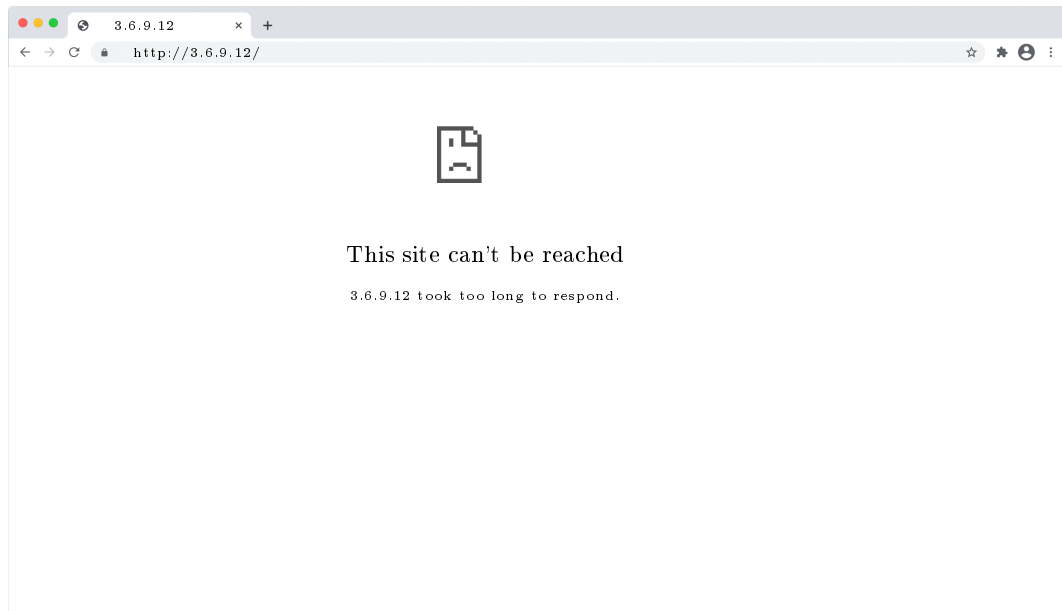And run the tests again

```
>> k6 run stress-test.js
```

Oh no...

```
1  status was 200
2  99% - 347867 / 130

4  checks.........................: 99.96%
5  data_received..................: 100 MB 44 kB/s
6  data_sent......................: 27 MB 12 kB/s
7  iterations.....................: 347997 152.552084/s
8  vus............................: 1 min=1 max=400
```

This site can't be reached

3.6.9.12 took too long to respond.

How can we fix this?

How can we fix this?

More servers?

```
» cat hello-server.tf
resource "aws_instance" "hello-server" {
    ami = "ami-04902260ca3d33422"
    instance_type = "t2.micro"

    user_data = file("./setup.sh")

    security_groups = [
        aws_security_group.hello-server.name
    ]

    tags = {
        Name = "hello-server"
    }
}
```

```
» cat hello-scale.tf
```

```
1  resource "aws_instance" "hello-server" {
2    count = 4
3
4    ami = "ami-04902260ca3d33422"
5    instance_type = "t2.micro"
6    user_data = file("${path.module}/setup.sh")
7
8    security_groups = [
9      aws_security_group.hello-server.name
10   ]
11
12   tags = {
13     Name = "hello-server-${count.index}"
14   }
15 }
```

Definition 2. Target Group

A collection of EC2 instances.
More specifically, a collection of network connection points to EC2 instances.

An empty HTTP target group

```
» cat hello-scale.tf

1  resource "aws_lb_target_group" "hello-target" {
2    name = "hello-target-group"
3    port = 80
4    protocol = "HTTP"
5    vpc_id = aws_security_group.hello-server.vpc_id
6  }
```

Definition 3. Health Check

Monitors attributes of hardware or software to detect deficiencies.

Add a health check

```
» cat hello-scale.tf

resource "aws_lb_target_group" "hello-target" {
  name = "hello-target-group"
  port = 80
  protocol = "HTTP"
  vpc_id = aws_security_group.hello-server.vpc_id

  health_check {
    port = 80
    protocol = "HTTP"
    timeout = 5
    interval = 10
  }
}
```

Add our instances to the target group

```
» cat hello-scale.tf

resource "aws_lb_target_group_attachment" "hello-target-link" {
  count = length(aws_instance.hello-server)
  target_group_arn = aws_lb_target_group.hello-target.arn
  target_id = aws_instance.hello-server[count.index].id
  port = 80
}
```

Definition 4. Load Balancer

A networking tool to route and distribute traffic to targets.

Create a load balancer

```
» cat hello-scale.tf
```

```
1  data "aws_subnet_ids" "nets" {
2      vpc_id = aws_security_group.hello-server.vpc_id
3  }
4
5  resource "aws_lb" "hello-balancer" {
6    name = "hello-balancer"
7    internal = false
8    load_balancer_type = "application"
9    subnets = aws_subnet_ids.nets.ids
10   security_groups = [
11     aws_security_group.hello-server.name
12   ]
13 }
```

Route load balancer traffic to the target group

```
» cat hello-scale.tf
```

```
1  resource "aws_lb_listener" "app" {
2    load_balancer_arn = aws_lb.hello-balancer.arn
3    port = "80"
4    protocol = "HTTP"

6    default_action {
7      type = "forward"
8      target_group_arn = aws_lb_target_group.hello-target.arn
9    }
10 }
```

We're live!

Hello, world! ×

http://hello-balancer-1992631166.us-east-1.elb.amazonaws.com/

Hello world from Brae

Use *k6* to determine the new *load limits*